*Last updated on* ***May 12, 2026***

# Generative AI models maintenance policy

This article describes the model maintenance policy for the Foundation Model APIs pay-per-token, Foundation Model APIs provisioned throughput, and Foundation Model Fine-tuning offerings.

In order to continue supporting the most state-of-the-art models, Databricks might update supported models or retire older models for these offerings.

## Model retirement policy

The model retirement policy explains how Databricks notifies you when a supported model is set for retirement, what happens during the transition period, and what to expect on the retirement date. Timelines differ by offering and model category as summarized in the following sections.

For currently retired models and planned retirement dates, see Retired models. For partner models, see Partner model retirement policy.

**IMPORTANT**

The retirement policies that apply to the Foundation Model APIs pay-per-token and Foundation Model Fine-tuning offerings only impact supported chat and completion models.

### Foundation Model APIs pay-per-token

The following table summarizes the retirement policy for Foundation Model APIs pay-per-token.

| Retirement notification | Transition to retirement | On the retirement date |
|---|---|---|
| Databricks takes the following steps to notify customers about a model that is set for retirement:<br>• On the **Serving** page of your Databricks workspace, a warning message appears on the model card that indicates that the model is planned for retirement.<br>• The applicable documentation contains a notice that indicates the model is planned for retirement and the start date it will no longer be supported. | Databricks will retire the model in **three months**. During this three-month period, customers can either:<br>• Choose to migrate to a Foundation Model APIs provisioned throughput endpoint to continue using the model past its end-of-life date.<br>• Migrate existing workflows to use recommended replacement models. | The model is no longer available for use and removed from the product. Applicable documentation is updated to recommend using a replacement model. |

# Foundation Model APIs provisioned throughput

The following table summarizes the retirement policy for Foundation Model APIs provisioned throughput.

| Retirement notification | Transition to retirement | On the retirement date |
|---|---|---|
| Databricks takes the following steps to notify | Databricks will retire the model in **six** | The model is no longer available for use and |

| Retirement notification | Transition to retirement | On the retirement date |
|---|---|---|
| customers about a model that is set for retirement:<br><br>• For endpoints that serve a deprecated model, a warning message appears on that serving endpoint's details page in your Databricks workspace. This message indicates that the model is planned for retirement and the applicable retirement date.<br>• A tooltip message provides recommended alternate models for workload migration.<br>• The applicable documentation contains a notice that indicates the model is planned for retirement and the start date it will no longer be supported. | **months**. During this six-month period:<br><br>• Customers can continue running existing provisioned throughput endpoints using the deprecated model until the retirement date.<br>• Customers that are not actively using a deprecated model cannot create new provisioned throughput endpoints or restart stopped endpoints for a deprecated model. | removed from the product.<br><br>• All endpoints using the retired model are transitioned to a failed state with a descriptive message. Any requests to these endpoints will fail.<br>• The customer can delete endpoints that use the retired model, but cannot restart them.<br>• Applicable documentation is updated to recommend using a replacement model. |

## Partner model retirement policy

Partner models are models provided by third-party partners, specifically OpenAI, Anthropic, and Google, that are available through Foundation Model APIs. For these partner models, Databricks generally follows the same deprecation timelines and policies as described for provisioned throughput and pay-per-token models.

However, retirement dates provided by partners might be shorter than the transition periods published by Databricks. In these cases, Databricks attempts to bridge the gap by temporarily redirecting models to a similar version, so customers receive the full transition time.

For example, if a pay-per-token model deprecation is announced with one month's lead time instead of three, Databricks redirects the model for an additional two months to prevent immediate breakage and allow time for migration. Queries fail at the end of the full three-month period.

> **NOTE**
>
> This redirection can only occur if the replacement model has the same price and is backwards compatible. The replacement model is usually an incremental model version, like 3.0 versus 3.1.

## Foundation Model Fine-tuning

The following table summarizes the retirement policy for Foundation Model Fine-tuning.

| Retirement notification | Transition to retirement | On the retirement date |
| --- | --- | --- |
| Databricks takes the following steps to notify customers about a model that is set for retirement:<br><br>• In the **Experiments** tab, a warning message appears in the dropdown menu for Foundation Model Fine-tuning that indicates that the model is planned for retirement. | Databricks retires the model in **three months**. During this three-month period, customers can migrate existing workflows to use recommended replacement models. | The model is no longer available for use and removed from the product. Applicable documentation is updated to recommend using a replacement model. |

| Retirement notification | Transition to retirement | On the retirement date |
|---|---|---|
| • The applicable documentation contains a notice that indicates the model is planned for retirement and the start date it will no longer be supported. | | |

# Model updates

Databricks might ship incremental model updates to deliver optimizations. When a model is updated, the endpoint URL remains the same, but the model ID in the response object changes to reflect the date of the update. For example, if an update is shipped to `meta-llama/Meta-Llama-3.3-70B` on 3/4/2024, the model name in the response object updates to `meta-llama/Meta-Llama-3.3-70B-030424`. Databricks maintains a version history of the updates that you can refer to.

# Retired models

The following sections summarize current and upcoming model retirements for the indicated feature offerings.

## Foundation Model APIs pay-per-token retirements

The following table shows model retirements, their retirement dates, and recommended replacement models to use for Foundation Model APIs pay-per-token serving workloads. Databricks recommends that you migrate your applications to use replacement models before the indicated retirement date.

| Partner model | Retirement date | Recommended replacement model |
| --- | --- | --- |
| Anthropic Claude 3.7 Sonnet | April 12, 2026 | Use the latest Claude Sonnet model |

| Open model | Retirement date | Recommended replacement model |
| --- | --- | --- |
| Meta Llama 3.1 405B | February 15, 2026 | OpenAI GPT OSS 120B |
| DBRX Instruct | April 30, 2025 | Meta-Llama-4-Maverick |
| Mixtral-8x7B Instruct | April 30, 2025 | Meta-Llama-4-Maverick |
| Meta-Llama-3.1-70B-Instruct | December 11, 2024 | Meta-Llama-4-Maverick |
| Meta-Llama-3-70B-Instruct | July 23, 2024 | Meta-Llama-4-Maverick |
| Meta-Llama-2-70B-Chat | October 30, 2024 | Meta-Llama-4-Maverick |
| MPT 7B Instruct | August 30, 2024 | Meta-Llama-4-Maverick |
| MPT 30B Instruct | August 30, 2024 | Meta-Llama-4-Maverick |

| Partner model | Retirement date | Recommended replacement model |
| --- | --- | --- |
| OpenAI GPT-5.1 Codex Max | July 16, 2026 | OpenAI GPT-5.3 Codex |

| Partner model | Retirement date | Recommended replacement model |
|---|---|---|
| OpenAI GPT-5.1 Codex Mini | July 16, 2026 | OpenAI GPT-5.3 Codex |
| OpenAI GPT-5.2 Codex | July 16, 2026 | OpenAI GPT-5.3 Codex |

If you require long-term support for a specific model version, Databricks recommends using Foundation Model APIs provisioned throughput for your serving workloads.

# Foundation Model APIs provisioned throughput retirements

The following table shows model family retirements, their retirement dates, and recommended replacement models to use for Foundation Model APIs provisioned throughput serving workloads. Databricks recommends that you migrate your applications to use replacement models before the indicated retirement date.

| Partner model | Retirement date | Recommended replacement model |
|---|---|---|
| Gemini 3 Pro | March 26, 2026 | Gemini 3.1 Pro. To allow more time for migration, between March 26, 2026 and June 7, 2026, API calls to Gemini 3 Pro will be temporarily redirected to Gemini 3.1 Pro. The pricing for both models is identical. |

| Open model family | Retirement date | Recommended replacement model |
|---|---|---|
| Meta Llama 3.1 405B | May 15, 2026 | OpenAI GPT OSS 120B |

| Open model family | Retirement date | Recommended replacement model |
| --- | --- | --- |
| Meta Llama 3 70B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| Meta Llama 3 8B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| Meta Llama 2 70B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| Meta Llama 2 13B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| Meta Llama 2 7B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| Mixtral 8x7B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| Mistral 7B | February 27, 2026 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| DBRX | December 19, 2025 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| MPT 30B | December 19, 2025 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |
| MPT 7B | December 19, 2025 | Comparable model on the same offering, like Llama 3.2, 3.3, or 4 model of similar size. |

## Foundation Model Fine-tuning retirements

The following table shows retired model families, their retirement dates, and recommended replacement model families to use for Foundation Model Fine-tuning workloads. Databricks recommends that you migrate your applications to use replacement models before the indicated retirement date.

| Model family | Retirement date | Recommended replacement model family |
|---|---|---|
| DBRX | April 30, 2025 | Llama-3.1-70B |
| Mixtral | April 30, 2025 | Llama-3.1-70B |
| Mistral | April 30, 2025 | Llama-3.1-8B |
| Meta-Llama-3.1-405B | January 30, 2025 | Llama-3.1-70B |
| Meta-Llama-3 | January 7, 2025 | Meta-Llama-3.1 |
| Meta-Llama-2 | January 7, 2025 | Meta-Llama-3.1 |
| Code Llama | January 7, 2025 | Meta-Llama-3.1 |

# Find workloads that use retired models

Use the following query to find workloads that are using deprecated models and identify their owners.

SQL

```sql
SELECT
   eu.requester,
   se.endpoint_name,
   se.entity_name,
   COUNT(*) AS request_count,
   SUM(eu.input_token_count) AS total_input_tokens,
   SUM(eu.output_token_count) AS total_output_tokens,
   MIN(eu.request_time) AS first_request,
```

Ask Genie

```sql
  MAX(eu.request_time) AS last_request
FROM system.serving.endpoint_usage eu
JOIN system.serving.served_entities se
  ON eu.served_entity_id = se.served_entity_id
WHERE LOWER(se.entity_name) LIKE '%claude-opus-4-1%'
GROUP BY eu.requester, se.endpoint_name, se.entity_name
ORDER BY request_count DESC
```