

Last updated on **Apr 6, 2026**

Share data and AI assets securely

Databricks provides tools to share data and AI assets securely with users outside of your organization or on different metastores within your account. Each feature serves different use cases:

- [Delta Sharing](#) for direct data sharing
- [Marketplace](#) for open data product distribution
- [Clean Rooms](#) for privacy-preserving collaboration

Delta Sharing

Delta Sharing enables you to share data and AI assets with users outside of your organization, regardless of whether they use Databricks. Share tabular and non-tabular data, AI models, views, filtered data, and notebooks using open protocols.

Topic	Description
Delta Sharing overview	Learn about Delta Sharing capabilities and use cases.
Set up Delta Sharing for your account	Configure Databricks as a provider to start sharing data with external users.
Databricks-to-Databricks sharing	Share data between Databricks accounts using secure, native integration.
Open sharing	Share data with users who don't use Databricks using open Delta Sharing protocols.

 Ask Genie

Topic	Description
Create and manage shares	Create shares to organize and control access to your data assets.
Access data shared with you	Access and query data that has been shared with your account as a recipient.
Use the SAP BDC connector for Databricks	Create a connection with SAP Business Data Cloud (BDC) to share data between SAP BDC and Databricks.

Marketplace

Marketplace is an open forum for exchanging data products. Share and discover datasets, notebooks, Solution Accelerators, and machine learning models.

Topic	Description
Databricks Marketplace overview	Learn about Marketplace capabilities and how to get started as a consumer or provider.
List data products (provider)	Publish your data products to Marketplace for others to discover and use.
Manage listings (provider)	Update and maintain your Marketplace listings.
Create private exchanges (provider)	Set up private exchanges to share data products with specific organizations.
Request and access data products (consumer)	Discover and access data products using Unity Catalog.
Manage requests and installed products (consumer)	Track and manage your data product requests and installations.

 Ask Genie

Clean Rooms

Clean Rooms provides a secure, privacy-protecting environment where multiple parties can collaborate on sensitive data without direct access to each other's underlying data.

Topic	Description
Clean Rooms overview	Learn about Clean Rooms capabilities and collaboration workflows.
Create clean rooms	Set up a clean room and invite collaborators to work with your data.
Work with output tables	Manage and access the results of clean room analyses.
Manage clean rooms	Learn how to update, monitor, and delete clean rooms.
Work with clean rooms as a collaborator	Accept invitations and collaborate in clean rooms created by other organizations.

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Apr 6, 2026**

What is Databricks Marketplace?

This article introduces Databricks Marketplace, an open forum for exchanging data products. Databricks Marketplace, which includes Partner Connect integrations, expands the Databricks ecosystem by enabling seamless integration with third-party tools and data services. Marketplace takes advantage of [Delta Sharing](#) to give data providers the tools to share data products securely and data consumers the power to explore and expand their access to the data and data services they need. Partner Connect further simplifies this by offering pre-configured integrations with select Databricks technology partners, making it easier for users to discover and connect to partner solutions.

The screenshot displays the Databricks Marketplace homepage. At the top, it says "Databricks Marketplace powered by DELTA SHARING". Navigation links for "Learn" and "My requests" are in the top right. The main heading is "Open marketplace for data, analytics, and AI", followed by a subtext: "Discover, evaluate, and gain access to Data sets, AI models, Solutions Accelerators, notebooks, and more...". Below this are buttons for "Browse free products" and "Suggest a product". A central diagram shows various partner logos (including AWS, Snowflake, and others) connected to a central Databricks logo. Below the diagram is a search bar and filters for "Product", "Provider", "Category", "Model task", "Free", and "Private exchange". The "Featured providers" section lists John Snow Labs, People Data Labs, Foursquare, S&P Global Commodity Insights, and IQVIA. The "Models" section lists several AI models from John Snow Labs, all marked as "Free", including "Clinical De-identification for Italian (Mask)", "Extract neurologic deficits related to Stroke Scale (NIHSS)", "Clinical De-identification for Arabic (Obfuscate)", "Clinical De-identification for Spanish (Obfuscate)", and "Extract Clinical Findings and the corresponding ICD 9 codes".

What kinds of data assets are shared on Databricks Marketplace?

Marketplace assets include datasets, Databricks notebooks, Databricks [Solution Accelerators](#), machine learning (AI) models, and Model Context Protocol (MCP) servers. Datasets are typically made available as catalogs of tabular data, although non-tabular data, in the form of Databricks [volumes](#), is also supported. Solution Accelerators are available as clonable Git repos. MCP servers provide tools and APIs that AI agents can use to interact with external systems.

What is Databricks Partner Connect?

Partner Connect lets you create trial accounts with select Databricks technology partners and connect your Databricks workspace to partner solutions from the Databricks UI. This allows you to try partner solutions using your data in the Databricks lakehouse, then adopt the solutions that best meet your business needs.

Partner Connect provides a simpler alternative to manual partner connections by provisioning the required Databricks resources on your behalf, then passing resource details to the partner. See [What is Databricks Partner Connect?](#).

How do consumers get access to data in Databricks Marketplace?

To find a data product you want on the Databricks Marketplace, simply browse or search provider listings.

You can browse:

- The [Open Marketplace](#), which does not require access to a Databricks workspace.
- The Databricks Marketplace on your Databricks workspace. Just click  **Marketplace**.

To request access to data products in the Marketplace, you must use the [Marketplace](#) on a Databricks workspace. You do not need a Databricks workspace to access and  Ask Genie

work with data once it is shared, although using a Databricks workspace with Unity Catalog enabled lets you take advantage of the deep integration of Unity Catalog with Delta Sharing.

Some data products are available to everyone in the *public marketplace*, and others are available as part of a *private exchange*, in which a provider shares their listings only with member consumers. Whether public or private, some data products are available instantly, as soon as you request them and agree to the terms. Others might require provider approval and transaction completion using provider interfaces. In either case, the Delta Sharing protocol that powers the Marketplace ensures that you can access shared data securely.

Get started accessing data products

To learn how to get started as a data consumer:

- Using a Databricks workspace that is enabled for Unity Catalog, see [Access data products in Databricks Marketplace \(Unity Catalog-enabled workspaces\)](#).
- Using third-party platforms like Power BI, pandas, or Apache Spark, along with Databricks workspaces that are not enabled for Unity Catalog, see [Access data products in Databricks Marketplace using external platforms](#).

To install and use MCP servers for AI agent development, see [Get access to external MCP servers](#).

How do providers list data products in Databricks Marketplace?

Databricks Marketplace gives data providers a secure platform for sharing data products that data scientists and analysts can use to help their organizations succeed. Databricks Marketplace uses Delta Sharing to provide security and control over your shared data. You can share public data, free sample data, and commercialized data offerings. You can share data products in public listings or as part of private exchanges that you create, making listings discoverable only by member consumers. In addition to datasets, you can also share Databricks notebooks, machine learning models, and

MCP servers to demonstrate use cases and provide tools that help customers succeed with AI agents.

Get started listing data products

To list your data products on Databricks Marketplace, you must:

- Have a Databricks account and premium workspace that is enabled for Unity Catalog. You do not need to enable all of your workspaces for Unity Catalog. You can create one specifically for managing Marketplace listings.
- Apply to be a provider through the [Databricks Data Partner Program](#). Alternatively, if you only want to share data through [private exchanges](#), you can use the self-service signup flow in the provider console. See [Sign up to be a Databricks Marketplace provider](#).
- Review the [Marketplace provider policies](#).

To learn how to get started, see [Become a Databricks Marketplace provider](#).

To learn how to create MCP server listings, see [MCP server listing fields](#).

View a demo

This video introduces Databricks Marketplace, shows how consumers access listings, and demonstrates how providers create them.

 **Databricks Marketplace**
Databricks

Databricks Marketplace

Open marketplace for data, analytics and AI

Steve Mahoney
Sr Manager, Product Management



©2022 Databricks Inc. — All rights reserved



Mirar en  YouTube

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Apr 6, 2026**

Access data products in Databricks Marketplace (Unity Catalog-enabled workspaces)

This article describes how to access data products in Databricks Marketplace if you have a Databricks workspace that is enabled for Unity Catalog.

NOTE

If you don't have a Databricks workspace that is enabled for Unity Catalog, you can access shared Marketplace data products using Delta Sharing open sharing connectors. See [Access data products in Databricks Marketplace using external platforms](#).

Overview

Databricks Marketplace gives you, as a data consumer, a secure platform for discovering data products that your organization needs to be successful. Databricks Marketplace uses Delta Sharing to provide security and control over shared data. Consumers can access public data, free sample data, and commercialized data offerings. Consumers who use a Unity Catalog-enabled Databricks workspace are not limited to accessing tabular data, but can also access volumes (non-tabular data), AI models, Databricks notebooks, and Databricks Solution Accelerators.

When you consume Marketplace data products using a Databricks workspace that is enabled for Unity Catalog, you can take advantage of the deep integration between Delta Sharing and Unity Catalog, along with Unity Catalog governance, auditing, and convenient interfaces.

Before you begin

To browse data product listings on Databricks Marketplace, you can use either of the following:

- The [Open Marketplace](#).
- A Databricks workspace.

To consume data products using a Databricks workspace that is enabled for Unity Catalog, you must have the following:

- A Databricks account on the [Premium plan or above](#).
- A Databricks workspace that is enabled for Unity Catalog (of course). See [Enable a workspace for Unity Catalog](#).

If you don't have one, you can get a free trial. Click **Try for free** on the [Open Marketplace](#) and follow the prompts to start your trial.

❗ IMPORTANT

You don't need to enable all of your workspaces for Unity Catalog. You can create a new one and enable it for Unity Catalog, using that workspace to receive Marketplace data products. If this option is not available, use the Marketplace on external platforms option. See [Access data products in Databricks Marketplace using external platforms](#).

To learn how to enable a workspace for Unity Catalog, see [Get started with Unity Catalog](#).

- `USE MARKETPLACE ASSETS` privilege on the Unity Catalog metastore attached to the workspace. See [Privileges that apply only to Delta Sharing or Databricks Marketplace](#). This privilege is enabled for all users on all Unity Catalog metastores by default.

If your admin has disabled this privilege, you can request that they grant it to you or that they grant you either of these:

- `CREATE CATALOG` and `USE PROVIDER` permissions on the Unity Catalog metastore.
- [Metastore admin role](#).

If you do not have any of these privileges, you can still view Marketplace listings but cannot access data products using Unity Catalog.

For more information, see [Unity Catalog privileges reference](#) and [Manage privileges in Unity Catalog](#). See also [Disable Marketplace access](#).

- If you want to install an MCP server, you must have the `CREATE CONNECTION` privilege on the Unity Catalog metastore to create the underlying connection object. Your workspace must also have the **Managed MCP Servers** preview enabled. See [Manage Databricks previews](#).

Browse Databricks Marketplace listings

To find a data product you want, simply browse or search the data product listings in Databricks Marketplace.

NOTE

As an alternative to the instructions that follow, you can search for Marketplace listings using the global search bar at the top of your Databricks workspace. See [Search for workspace objects](#). You can also view and request free sample data on the **Add data** page. In the workspace sidebar, click **Data ingestion** and scroll down to **Free sample data from Databricks Marketplace**.

1. Go to marketplace.databricks.com or log into your Databricks workspace and click  **Marketplace**.
2. Browse or search for the data product that you want.

Click the  filter icon to filter listings by product type (dataset, [solution accelerator](#), or ML model), provider name, category, cost (free or paid), and more.

If you are logged into a Databricks workspace, you can also choose to view only the private listings available to you as part of a private exchange. See [Participate in private exchanges](#).

Request access to data products in the Marketplace

To request access to data products, you must be logged into a Databricks workspace. Some data products are available immediately, and others require provider approval and transaction completion using provider interfaces.

Requirements

See [Before you begin](#). To access data products in the Marketplace, you must have at least the `USE MARKETPLACE ASSETS` privilege on the Unity Catalog metastore that is attached to the workspace that you are using.

Get access to data products that are instantly available

Some data products are available instantly, requiring only that you request them and agree to terms. These are listed under the **Free and instantly available** heading on the Marketplace landing page, are identified on the listing tile as **Free**, and are identified as **Instantly available** on the listing detail page.

1. When you've found a listing you're interested in on the Marketplace landing page, click the listing to open the listing detail page.
2. Click the **Get instant access** button and accept the Databricks terms and conditions.

Accessing Databricks Solution Accelerators works a little differently. See [Get access to Databricks Solution Accelerators](#).

3. (Optional) Under **More options**, modify the suggested catalog name.

The catalog name is displayed in Catalog Explorer in your Databricks workspace, and it is used in the three-part namespace (catalog.schema.table|volume|view) in queries. You can change the default name.

4. Click the **Get instant access** button.
5. Click the **Open** button to view the data product, which appears as a read-only catalog in Catalog Explorer.

For information about managing access to catalogs in Databricks, see [Access the shared data using Unity Catalog](#).

If sample notebooks are available, they appear under the **Sample notebook** heading in the listing.

To view a notebook, click the **Preview notebook** button. To import a notebook into your Databricks workspace so that you can run it, click **Preview notebook** and then click **Import Notebook**. See [Databricks notebooks](#).

Request data products that require provider approval

Some data products require provider approval, typically because a commercial transaction is involved, or the provider might prefer to customize data products for you. These listings are identified on the listing detail page as **By request** and include a **Request access** button.

1. When you've found a listing you're interested in on the Marketplace landing page, click the listing to open the listing detail page.
2. Click the **Request access** button.
3. Enter your name, company, and a brief description of your intended use for the data product.
4. Accept the Databricks terms and conditions and click **Request access**.
5. You will be notified by email when the provider has completed their review of your request.

You can also monitor the progress of your request on the My Requests page in Marketplace. See [Manage shared Databricks Marketplace data products](#). However, any transactions that follow will use provider communications and payment platforms. No commercial transactions are handled directly on Databricks Marketplace.

6. When your transaction is complete, the data provider will make the data product available to you as a read-only catalog in your workspace.

For information about managing access to catalogs in Databricks, see [Access the shared data using Unity Catalog](#).

Get access to Databricks Solution Accelerators

Unlike other data assets, Databricks [Solution Accelerators](#) are shared by cloning Git repositories and making them available in [Databricks Git folders](#). To access a Solution Accelerator:

1. In Marketplace, click the Solution Accelerator listing to open the listing detail page.
2. Click the **Get instant access** button.
3. On the **Add Git folder** dialog, enter a name for the repository. This name will appear in Databricks Git folders UIs.
4. Accept the Databricks terms and conditions and click **Create Git folder**.
5. Click the **Open** button and select **Git folder** to view the repo in the workspace file browser.

Get access to external MCP servers

ⓘ PREVIEW

This feature is in [Public Preview](#).

Model Context Protocol (MCP) servers provide tools and APIs that AI agents can use to interact with external systems. In Databricks Marketplace, you can install external MCP servers to use with AI agent development in Databricks, including [Supervisor Agent](#) and [AI Playground](#).

When you install an external MCP server from Marketplace, it creates a Unity Catalog connection that your AI agents can use to access these tools. You can create multiple connections to the same MCP server.

For more information on using installed MCP servers in your agents or connecting to custom MCP servers not in Marketplace, see [Install an external MCP server](#).

To install an MCP server:

1. Find the MCP server listing in the Marketplace or in the **Agents > MCP Servers** tab.
2. Click **Install** on the listing card or detail page.
3. On the installation dialog, configure the connection:
 - **Connection name:** Enter a name for the Unity Catalog connection.
 - **Host:** The host domain for this MCP server. This field might be pre-populated.
 - **Base path:** The base path for this MCP server. This field might be pre-populated.

If the credential type is **Bearer token**, enter the following credential:

- **Bearer Token:** Provide your bearer token.

If the credential type is **OAuth U2M**, enter the following credentials:

- **Client ID:** Your client ID.
- **Client secret:** Your client secret.

Refer to the MCP server provider's documentation for instructions on obtaining credentials.

4. Click **Install**

After installation completes, the page displays the connection details, where you can view connection properties and grant other users access to the MCP server.

Access the shared data using Unity Catalog

Access to the data in your Databricks workspace depends on the data product type:

- Marketplace datasets and models are shared with you in a read-only catalog in Databricks. Catalogs are the top-level container for data assets that are managed by Unity Catalog. For more information about the data object hierarchy in Unity Catalog, see [Database objects in Databricks](#).

Once the provider has shared the data with you, you have a few ways of accessing the catalog. The sections that follow describe some of these access options.

- Notebooks are shared directly in the Marketplace listing, and you can import them from the listing into your workspace.
- Solution Accelerators are shared as Databricks Git folders. See [Get access to Databricks Solution Accelerators](#).

Access shared datasets in the Marketplace

To access shared data from the Marketplace:

1. In the sidebar, click  **Marketplace**.
2. On the Marketplace landing page, click **My requests** in the upper right-hand corner.
3. On the **Installed data products** tab, find the data product, click the  kebab menu at the end of the data product row, and select **View data**.

You can also click the data product name to open the data product details page, where you can click the **Open** button to view the data.

Catalog Explorer opens to the catalog that contains the data set, where you can access the data or manage access for other team members. See [Grant access to other team members](#).

To learn more about accessing shared data in Databricks, see [Read data shared using Databricks-to-Databricks Delta Sharing \(for recipients\)](#).

Access shared datasets in Catalog Explorer

To access shared data directly from Catalog Explorer:

1. In the sidebar, click  **Catalog**.
2. If you know the name of the catalog that holds the shared data, search for it and select it.

If you don't know the catalog name but you do know the provider name, you can find the catalog in Catalog Explorer by doing the following:

1. In the sidebar, click  **Catalog**.
2. In the left pane of Catalog Explorer, click **Delta Sharing** and then **Shared with me**.
3. On the **Providers** tab, click the provider name.
4. On the **Shares** tab, find the catalog name and click it to open it.

Access shared datasets using the Databricks CLI or SQL statements

You can also find and access the catalog that contains the shared data using the Databricks CLI or SQL statements in a Databricks notebook or Databricks SQL editor query. For details, see [Access data in a shared table or volume](#). You can skip the sections that describe how to create a catalog, since Databricks Marketplace does this for you.

Grant access to other team members

If you are the user who requested the shared data, you are the owner of the catalog that contains that data in your workspace. As such, you can grant your team members access to the catalog and refine access at the schema, table, view, row, and column level, just as you do any data in Unity Catalog. See [Manage privileges in Unity Catalog](#). That said, table and view data under a shared catalog is read-only, which means that you can only grant your team read operations like `DESCRIBE`, `SHOW`, and `SELECT`.  Ask Genie

You can also transfer ownership of the catalog or the objects inside it.

View sample notebooks

Some listings include sample notebooks on the listing details page. To access these notebooks for instantly available listings that have been shared with you:

1. In the sidebar, click  **Marketplace**.
2. On the Marketplace landing page, click **My requests** in the upper right-hand corner.
3. On the **Installed data products** tab, find the data product and click the data product name to open the listing details page. If there is a sample notebook, it appears under the **Sample notebook** heading in the listing.
4. Click **Preview notebook** to view the notebook, and click **Import notebook** to import it to your Databricks workspace.

NOTE

The **Sample notebooks** display and preview in the listings UI does not work in Chrome Incognito mode.

Participate in private exchanges

Some data providers might want to share certain data with a limited set of consumers who are invited to be part of a *private exchange*. You can find private exchange listings by selecting the **Private exchange** checkbox on the Marketplace home page. Just as you can with public listings, you can access free listings instantly or request access to those that are marked **By request**.

To join a private exchange, a data provider needs to invite your organization. When they do, they will request a shared identifier for your Unity Catalog metastore. To learn how to get your metastore's sharing identifier, see step 1 in [Get access in the Databricks-to-Databricks model](#).

To learn more about private exchanges, see [Create and manage private exchanges in Databricks Marketplace](#).

Disable Marketplace access

By default, all users in any Unity Catalog-enabled workspace have the ability to request data products in Databricks Marketplace. In other words, the `account users` group has the `USE MARKETPLACE ASSETS` privilege on all Unity Catalog metastores unless a metastore admin revokes that privilege.

This privilege does not grant the ability to participate in financial transactions with data providers. All financial transactions take place outside Databricks. This privilege does grant the ability to access data products that are labelled **Free and instantly available** and data products that have already been purchased. Accessing such data products creates new read-only catalogs in Databricks that are owned by the requestor, who can grant read-only access to other users.

A metastore admin can disable all users' ability to request data products in Databricks Marketplace by revoking the `USE MARKETPLACE ASSETS` privilege from the `account users` group on the Unity Catalog metastore. If you do revoke this privilege, users can continue to browse the Databricks Marketplace in their workspace but cannot request data products.

Permission required: Metastore admin

NOTE

If your workspace was enabled for Unity Catalog automatically, you might not have a metastore admin. For more information, see [Automatic enablement of Unity Catalog](#).

Catalog Explorer **SQL**

To revoke the `USE MARKETPLACE ASSETS` privilege for the `account users` group:

1. In your Databricks workspace, click  **Catalog**.
2. Click the link icon next to the metastore name at the top left.
3. Find the row that grants `account users` the `USE MARKETPLACE ASSETS` privilege.
4. Click the checkbox next to the row and click the **Revoke** button.
5. Confirm the revoke action.

 Ask Genie

6. Grant the privilege to any specific users and groups you like by clicking the **Grant button**.

If you don't want your users to be able to view the Marketplace home page at all, contact your Databricks account team.

Known issues

If a request for access is rejected by the data provider, you cannot request the same data product again. If you encounter this issue, contact your provider or Databricks account team.

Share your feedback

We'd love to hear what you think about Databricks Marketplace.

- Have feedback? Go to **Learn > Provide feedback** on the Marketplace home page.
- Want to see additional datasets in the marketplace? Click **Suggest a product** on the Marketplace home page banner.

Next steps

- [Manage Marketplace requests and data that has been shared with you](#)

Is this page

as this page helpful?

☐ Yes

☐ No

 Ask Genie

Last updated on **Jan 20, 2026**

Access data products in Databricks Marketplace using external platforms

This article describes how to access data products in Databricks Marketplace without a Unity Catalog-enabled Databricks workspace. You can use Delta Sharing open sharing connectors to access Marketplace data using a number of common platforms, including Microsoft Power BI, Microsoft Excel, pandas, Apache Spark, and non-Unity Catalog Databricks workspaces. Only tabular data sets are available on external platforms (not Databricks notebooks, volumes, or models, for example).

❗ IMPORTANT

If you have a Databricks workspace that is enabled for Unity Catalog, you should access shared data using Unity Catalog. See [Access data products in Databricks Marketplace \(Unity Catalog-enabled workspaces\)](#).

Before you begin

To browse data product listings on Databricks Marketplace, you can use either of the following:

- The [Open Marketplace](#).
- A Databricks workspace.

To request access to data products, regardless of platform, you must have a Databricks workspace.

If you don't have one, you can get a free trial. Click **Try for free** on the [Open Marketplace](#) and follow the prompts to start your trial.

Browse Databricks Marketplace listings that are accessible on external platforms

To find a data product you want, browse or search the data product listings in Databricks Marketplace. Only **Data set** (tabular data) product types are available to share using external platforms or legacy Databricks workspaces that don't use Unity Catalog.

1. Go to marketplace.databricks.com or log into your Databricks workspace and click  **Marketplace**.

Alternatively, you can search for Marketplace listings using the global search bar at the top of your Databricks workspace. See [Search for workspace objects](#).

1. Browse or search for the data product that you want.

You can filter listings by provider name, product type, category, cost (free or paid), or keyword search. Only the **Data set** product type is available for access using external platforms.

If you are logged into a Databricks workspace, you can also choose to view only the private listings available to you as part of a private exchange. See [Participate in private exchanges](#).

Get access to data products that are accessible on external platforms

To request access to data products, you must be logged into a Databricks workspace, even if you will use the shared data product on an external platform. Some data products are available immediately, and others require provider approval and transaction completion using provider interfaces.

 Ask Genie

Get access to data products that are instantly available

Data products that are available instantly require only that you request them and agree to the terms. These data products are listed under the **Free and instantly available** heading on the Marketplace landing page, are identified on the listing tile as **Free**, and are identified as **Instantly available** on the listing detail page.

1. When you've found a listing you're interested in on the Marketplace landing page, click the listing to open the listing detail page.
2. Click the **Get instant access** button.
3. Under **More options** , select **On external platforms**.
4. Accept the Databricks terms and conditions.
5. Click **Get instant access**.
6. Click the **Download credential file** button to get the credential file, which you and your team can use to gain access to shared data using third-party data platforms and non-Unity Catalog Databricks workspaces.

❗ IMPORTANT

The credential file can only be downloaded once. The download button remains active after you download the file, but subsequent downloads rotate to a new credential. The old credential expires after one day or its original expiration date, whichever is sooner. Only two credentials can be active at the same time.

7. Store the credential file in a secure location.

Don't share the credential file with anyone outside the group of users who should have access to the shared data. If you need to share it with someone in your organization, Databricks recommends using a password manager.

To learn how to access the shared data using your platform of choice, see [Access shared data using Delta Sharing open sharing connectors](#).

 Ask Genie

Request data products that require provider approval

Some data products require provider approval, typically because a commercial transaction is involved or the provider prefers to customize data products for you. These listings are identified on the listing detail page as **By request** and include a **Request access** button.

1. When you've found a listing you're interested in on the Marketplace landing page, click the listing to open the listing detail page.
2. Click the **Request access** button.
3. Enter your name, company, and a brief description of your intended use for the data product.
4. Click **More options** and select **On external platforms**.
5. Accept the Databricks terms and conditions and click **Request access**.
6. You will be notified by email when the provider has completed their review of your request.

You can also monitor the progress of your request on the My Requests page in Marketplace. See [Manage shared Databricks Marketplace data products](#). However, any transactions that follow will use provider communications and payment platforms. No commercial transactions are handled directly on Databricks Marketplace.

7. When your transaction is complete, you will receive a notification email from the data provider, and the listing will display a **Download credential file** button. Click this button to download the credential file, which you and your team can use to gain access to shared data using third-party data platforms and non-Unity Catalog Databricks workspaces.

You can also find the listing in Marketplace under **My requests**. When the credential is ready for download, the data product appears on the **Installed data products** tab.

❗ **IMPORTANT**

The credential file can **only be downloaded once**. The download button remains active after you download the file, but subsequent downloads rotate to a new credential. The old credential expires after one day or its original expiration date, whichever is sooner. Only two credentials can be active at the same time.

8. Store the credential file in a secure location.

Don't share the credential file with anyone outside the group of users who should have access to the shared data. If you need to share it with someone in your organization, Databricks recommends using a password manager.

To learn how to access the shared data using your platform of choice, see [Access shared data using Delta Sharing open sharing connectors](#).

Access shared data using Delta Sharing open sharing connectors

To use external platforms or non-Unity-Catalog Databricks workspaces to access datasets that have been shared using Databricks Marketplace, you need the credential file that was downloaded from the Marketplace listing. You use this credential file to access the shared data using Delta Sharing open sharing connectors.

For full instructions for using non-Unity-Catalog Databricks workspaces, Apache Spark, pandas, and Power BI to access and read shared data, see [Read data shared using Delta Sharing open sharing with bearer tokens](#).

For a full list of Delta Sharing connectors and information about how to use them, see the [Delta Sharing open source documentation](#).

Limitations on sharing to external platforms using Marketplace

Some tables require partition information from the consumer side (for example, partitioning by country or region). In the open sharing protocol, this information is not available to the share, and therefore, the table is not accessible. The following error is returned: "Recipient authentication failure: the data is restricted by Recipient properties that do not apply to the current recipient in the session. Please contact the data provider to resolve the issue."

Is this page

as this page helpful?

☐ Yes

☐ No

Last updated on **Mar 26, 2026**

Manage shared Databricks Marketplace data products

This article describes how to manage data products that are shared with you through Databricks Marketplace and your pending requests for access to data products. This article is intended for data consumers.

To learn how to request and access data products in Databricks Marketplace, see [Access data products in Databricks Marketplace \(Unity Catalog-enabled workspaces\)](#).

Before you begin

To manage data products that have been shared with you, you must have the `CREATE CATALOG` and `USE PROVIDER` permissions on the Unity Catalog metastore attached to your workspace, or the [metastore admin role](#).

To get these privileges, ask your Databricks account admin or metastore admin to grant them to you. For more information, see [Unity Catalog privileges reference](#).

If your workspace was enabled for Unity Catalog automatically, the workspace admin has these permissions and can grant them to other users. See [Automatic enablement of Unity Catalog](#).

View request status

If you have requested a data product that requires provider approval, you can view request status in the Marketplace UI:

1. Log into your Databricks workspace.

For required permissions, see [Before you begin](#).

2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **My requests**.
4. On the **Requests** tab, view all requests and their current review status: **Pending**, **Fulfilled**, and **Denied**, along with the requested and reviewed dates.

Transactions that are ongoing between you and a data provider occur outside of the Databricks Marketplace system. Details of those transactions are not captured here.

View and access installed data products

To view installed data products and access them in Catalog Explorer:

1. Log into your Databricks workspace.

For required permissions, see [Before you begin](#).

2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **My requests**.
4. On the **Installed data products** page, you can view all installed data products.

To view details about a specific product, such as provider info, documentation, terms, and sample notebooks, click the product name.

To view the shared data in Catalog Explorer, click the  kebab menu at the end of the data product row, and select **View data**.

Work with installed data products

Each Marketplace dataset is shared with you in a read-only catalog. Catalogs are the top-level container for data managed by Unity Catalog in Databricks.

 Ask Genie

For detailed information about accessing and granting permissions on shared catalogs, see [Access the shared data using Unity Catalog](#).

To learn how to access shared notebooks, see [View sample notebooks](#)

For more information about the data object hierarchy in Unity Catalog, see [Database objects in Databricks](#).

Delete installed data products

When you remove an installed data product, the shared catalog is deleted from your workspace. Remember to update or delete any queries, dashboards, or notebooks that depend on the data in that catalog.

To remove an installed data product:

1. Log into your Databricks workspace.

For required permissions, see [Before you begin](#).

2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **My requests**.
4. On the **Installed data products** tab, find the data product you want to delete and click the product name.
5. On the product detail page, click the  kebab menu next to the **Open** button and select **Uninstall product**.
6. On the confirmation dialog, type the repository's name (which is displayed in the sentence immediately above the confirmation field).
7. Click **Confirm and uninstall**.

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Apr 9, 2026**

Become a Databricks Marketplace provider

This page describes how to sign up as a Databricks Marketplace provider, assign the Marketplace admin role, and create your provider profile. After these steps are completed, you can create listings on Databricks Marketplace.

Before you begin

To list products in the Databricks Marketplace, you must agree to provider policies, and your account and workspaces must meet certain requirements.

Policies

To list data products on the Marketplace exchange, you must agree to the [Marketplace provider policies](#).

Account and workspace requirements

Databricks Marketplace uses [Delta Sharing](#) to manage secure sharing of data products. Delta Sharing, in turn, requires that your Databricks workspace is enabled for Unity Catalog. Your Databricks account and workspaces must therefore meet the following requirements:

- Databricks account on the [Premium plan or above](#).
- A Databricks workspace enabled for Unity Catalog. See [Enable a workspace for Unity Catalog](#). You do not need to enable all of your workspaces for Unity Catalog. You can create one to manage Marketplace listings.

If you have workspaces that meet these criteria, your users will be able to view the Marketplace home page. Additional permissions are required to create and manage listings. These are enumerated in the sections that follow. If you don't want your users to be able to view the Marketplace home page at all, contact your Databricks account team.

Permission requirements

To sign up as a private-exchange-only provider, you must be a Databricks account admin. See [Sign up to be a Databricks Marketplace provider](#).

Sign up to be a Databricks Marketplace provider

The way you sign up to be a Databricks Marketplace provider depends on whether you intend to create listings in the public marketplace or only through private exchanges. In private exchanges, only consumers who are members of the exchange can browse, view, and request access to a listing. See [Create and manage private exchanges in Databricks Marketplace](#).

To be a private-exchange-only provider, you can sign up using the provider console. To create public listings, you apply through the Databricks Data Partner Program.

Apply to be a provider that can create public listings

NOTE

If your organization is already in the Databricks Partner Program and you're interested in becoming a Marketplace provider, skip the following instructions and contact partnerops@databricks.com instead.

1. On the [Databricks Data Partner Program page](#), click **Apply Now**.
2. On the next page, click **Apply here** as a new provider.

3. Enter your email address and click **Apply Now**.
4. Fill out the application form.

Toward the bottom of the application form, you're asked which Databricks Partner program interests you. Select **Databricks Data Provider**.

The Databricks Partner team reaches out to you to complete the application process. When you're approved, the Provider console becomes available in your Unity Catalog-enabled Databricks workspaces. To access the Provider console, a user must have the Marketplace admin role.

Sign up as a private exchange provider

To sign up as a private-exchange-only provider, use the **Get started as a provider** page in the provider console.



Get started as a provider

Before you can start publishing on Databricks Marketplace, follow the steps below.

1

Assign Marketplace admin
Become a Marketplace admin. Marketplace admins can create and publish listings. Request your account admin to assign this role to you in the account console.

2

Create a provider profile
The profile will be the organization under which you will publish listings.

 **Assign Marketplace admin**

 **Refresh page**

To sign up as a private exchange provider:

1. As a Databricks account admin, sign in to your Databricks workspace.
2. In the sidebar, click  **Marketplace**.

3. In the upper-right corner of the Marketplace page, click **Provider console**.
4. If your account isn't onboarded as a provider, the provider console displays the **Get started as a provider** page. Follow the steps to enable your account as a private exchange provider.
5. After a few minutes, the **Assign Marketplace admin** button appears. Click it to open your user page in the account console. On the **Roles** tab, enable **Marketplace admin**.

You can assign the Marketplace admin role to other users. Those users can continue the process after you finish signing up as a private exchange provider.

6. Return to the **Provider console** in the workspace and click **Refresh page**. Don't use your browser's refresh button.
7. After a few minutes, the **Create provider profile** button appears. Click it to open the **Create Profile** page.

NOTE

It might take a few minutes for the system to finish assigning the Marketplace admin role. If you proceed to create a provider profile and see an error indicating that you don't have the Marketplace admin role, wait a few more minutes, refresh the page, and retry.

8. To create your provider profile, follow the instructions in [Create your Marketplace provider profile](#), starting at step 5.
9. Create your first private exchange. See [Create and manage private exchanges in Databricks Marketplace](#).

Assign the Marketplace admin role

If you signed up as a private exchange provider, you performed this task as part of the sign-up process. You can skip the instructions in this section unless you want to enable other users in your Databricks account as Marketplace admins.

After you've been approved as a Marketplace provider, you must grant at least one user the Marketplace admin role. This role is required for accessing the Marketplace Provider console and for creating and managing your Marketplace provider profile and listings. A Databricks account admin can grant the role.

1. As an account admin, log in to the [account console](#).
2. Click  **User management**.
3. Find and click the username.
4. On the **Roles** tab, turn on **Marketplace admin**.

Create your Marketplace provider profile

Your provider profile gives you the opportunity to tell prospective consumers who you are and to group your data products under a single brand or identity. Typically, a data provider has one profile but can list multiple data products. If you want more than one profile, reach out to your Databricks account team.

If you are a public provider, you can create your profile after your provider application has been approved. If you are a private-exchange-only provider, you create your profile as the final step of the [signup process](#), starting at step 5 in the procedure provided here.

Permission required: Marketplace admin role

To create a profile:

1. Log in to the Databricks workspace you will use for creating shares and listings.
2. In the sidebar, click  **Marketplace**.
3. In the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Provider console** page **Profiles** tab, click **Create profile**.
5. Enter the following information. All fields are required:
 - **Provider name:** Use a name that consumers will recognize. Consumers can filter listings by provider name.
 - **Logo:** Drag and drop or browse to an image file for the logo you want to use in your listing. You can upload different logos for light and dark mode.

- **Description:** Describe your organization clearly and accurately. Include details such as industries you typically serve or represent and the types of data assets that you typically list. Consumers can see this description when they view your profile and on all of your listings.
- **Organization website:** Link to your organization's website. Consumers can follow this link to learn more about your organization. This link appears on all of your listings.
- **Business email:** Enter an email address that Databricks can use to send you notifications. Consumers do not see this information.
- **Support email:** Enter an email address that consumers can use to request support. This address appears on all of your listings.
- **Terms of service link:** This link appears on all of your listings. You can override this link by entering a different one at the listing level.
- **Privacy policy:** This link appears on all of your listings. You can override this link by entering a different one at the listing level.

6. Click **Submit profile**.

If any of this information changes while you are a Marketplace provider, update your profile. Your profile must be accurate and kept up to date.

Next steps

- [Create a Marketplace listing](#)
- [Create and manage private exchanges in Databricks Marketplace](#)

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

 Ask Genie

Last updated on **Apr 6, 2026**

Create a Marketplace listing

This page describes how to create shares and Databricks Marketplace listings for your data products.

Prerequisites

You or your Databricks administrator has completed the Databricks Marketplace provider setup process:

- Sign up to become a Databricks Marketplace provider, public or private
- Create a Marketplace provider profile
- Have the Marketplace admin role assigned to you

For more information, see [Become a Databricks Marketplace provider](#).

Create shares

NOTE

To list a data product that is free and instantly available to consumers, you must include a share when you create the listing. Listings that require you to approve a consumer request, on the other hand, don't require that you include a share in the listing. You can create the share later, after any business agreements are complete and you've approved the consumer's request. If that's what you want to do, skip ahead to [Create a listing](#).

After you have a Databricks account enabled for Delta Sharing and a Databricks workspace enabled for Unity Catalog, you can create the *shares* that you use to share your data in the Marketplace.

A share is a Delta Sharing object. It's a collection of tables, views, volumes, and AI models that are shareable and securable as a unit. Tables can be shared with any consumer. Volumes, AI models, and notebooks can only be shared with consumers who have access to a Databricks workspace that is enabled for Unity Catalog.

1. Add data tables, views, or volumes to your Unity Catalog metastore.

To learn how to create these data assets in Unity Catalog, see:

- [Databricks tables](#)
- [Create and manage views](#)
- [What are Unity Catalog volumes?](#)

2. Create a share and add these data assets to the share.

To learn how to create and update shares, including permissions required, see [Create and manage shares for Delta Sharing](#).

After your share is created, you can create or update a Marketplace listing that references it.

Notebook example: Sample notebook

In addition to tables, volumes, and views, Databricks recommends sharing Databricks notebooks. Notebooks can help demonstrate example use cases and visualize table properties. Your listing can include sample notebook previews that consumers can import into their workspaces.

Crude Oil Prices: West Texas Intermediate - Cushing, Oklahoma | FRED Preview



Overview

This release contains information on West Texas Intermediate (WTI) crude oil prices. WTI is usually delivered in Cushing, Oklahoma. Cushing, Oklahoma is an important location for oil due to world oil prices being determined there as a result of WTIs being the most traded oil futures contract globally.

Product Details

Includes the following tables:

- **fs_fred.fred_DCOILWTICO**
 - **Update frequency:** Daily
 - **Source:** FRED
 - **Description:** The included dataset contains the following columns: DATE, DCOILWTICO. The DCOILWTICO column represents dollars per barrel at the time of the corresponding DATE.

Sample notebook

Crude Oil Prices: West Texas Intermediate - Cushing, Oklahoma | FRED

Cushing, OK: important location for oil due to world oil prices being determined there as a result of WTIs being the most traded oil futures contract globally.

[Preview notebook](#)

Tables Available in this Share

```
share_name = "fs_fred_dcoilwtico"
df_tables = spark.sql(f"SHOW ALL IN SHARE fs_fred_dcoilwtico")
print(f"This share contains 1 table(s)")
```

This share contains 1 table(s)

For more information about creating notebooks, see [Databricks notebooks](#). If you need help creating an effective sample notebook, contact dataprovers@databricks.com.

NOTE

The **Sample notebooks** display and preview in the listings UI do not work in Chrome Incognito mode.

The following example shows how to create an effective sample notebook. It includes guidance for creating sample notebooks for your listings.

Marketplace starter notebook for data providers

[Open notebook in new tab](#)

 [Copy link for import](#)

Marketplace Starter Notebook for Data Providers

The goal of this notebook is to provide an outline for a common workflow that consumers can use to examine new data shared with them. This is not an exhaustive data exploration presentation. When you create sample notebooks for your own listings, you are encouraged to highlight interesting examples to demonstrate the value of the data.

Expand notebook ▾

When you build your notebooks, it's often good to show:

Create a listing

Marketplace listings enable consumers to browse, select, and access your data products. All dataset listings are automatically shareable with both consumers on Databricks workspaces and consumers on third-party platforms like Power BI, pandas, and Apache Spark.

NOTE

Some data assets, like Unity Catalog volumes, can only be shared with consumers who have access to a Unity Catalog-enabled Databricks workspace. Tables, however, are shareable with all consumers. If you include both tables and volumes in a share, consumers who don't have access to a Unity Catalog-enabled workspace can only access the tabular data.

Permissions required: [Marketplace admin role](#). If you are creating and managing personalized listings (those that require provider approval before they're fulfilled), you must also have the `CREATE RECIPIENT` and `USE RECIPIENT` privileges. See [Unity Catalog privileges reference](#).

To create a listing:

 Ask Genie

1. Log in to your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. In the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Provider console** page **Listings** tab, click **Create listing**.
5. On the **New listing** page, enter your listing information. Each page in the listing wizard is described in [Fill out listing fields and options](#).

Fill out listing fields and options

The listing wizard guides you through publishing a Marketplace listing. This section describes the fields and options on each page, along with recommendations for creating an effective listing. During each step, you can save your progress as a draft and view a preview before you publish.

After you fill out all the fields, click **Publish**.

Step 1: General

Specify the basic information about your listing, such as the name and the provider profile to publish the listing under.

Field name	Description
Listing name	Each listing should have a unique name that helps consumers understand what it offers. Recommendations: • Fewer than 100 characters. • Title case (capitalize primary words).
Provider profile	Your organization or company name. Select from the drop-down menu. Databricks creates your profile as part of the partner organization approval process.
 Ask Genie	

Field name	Description
How can consumers access data assets in the listing?	This field cannot be changed after the listing is saved as a draft or published. Choose how your consumers gain access to the listing: • Data assets are instantly accessible: Allow consumers to gain access to the shared data directly from the Marketplace, without requiring your approval (but requiring acceptance of terms of service). • Consumers must request access to the listing: Select this option to require your approval before a consumer can access the shared data. Use this option if you require a business agreement before you make the data product available to consumers. You must manage the business agreement with consumers outside of Databricks Marketplace. You can initiate communications using the consumer email address. View and handle consumer requests on the Provider Console > Consumer requests tab. See Manage requests for your data product in Databricks Marketplace.
Publish to	Choose one of the following: • Public Marketplace ◦ All consumers can browse and view the listing in the public Databricks Marketplace. • Private exchange ◦ Only consumers who are members of a private exchange, created by you or another Marketplace admin, can browse, view, and request the listing. See Create and manage private exchanges in Databricks Marketplace. You must select at least one private exchange from the drop-down list.

Step 2: Data assets

Select which types of data assets to include in your listing.

Field name	Description
Tables	Include structured tabular data from your share.
Files	Volumes that contain a variety of data, such as images, audio, video, and more.
Models	AI models that consumers can deploy for inference.
Notebooks	Interactive notebooks for data science and machine learning workflows. For more information, see Notebook example: Sample notebook .
Solution accelerators	Prebuilt solution templates hosted in a Git repository that consumers can clone.
MCP server (Public Preview)	Model Context Protocol (MCP) servers are tools and APIs that AI agents can use to interact with external systems. Only other users using Databricks can install MCP servers listed on the Marketplace. For more details on the setup process, see MCP server listing fields. When you select this option, fill in the required fields, and advance to the next step, the listing is saved automatically.
Select share	If you selected to include tables, files, models, or notebooks in your listing, select the share that contains the data assets for your listing. You can create a new share if needed. For more information, see Create shares .

Step 3: Attributes

Add attributes and categories to make your listing more discoverable.

Field name	Description
Categories	Select the categories that consumers can use to filter listings. Categories also appear as tags on listing tiles and detail pages.
Update frequency (optional)	Specify how often the data asset in your listing is updated on Databricks.
Advanced attributes	Advanced data attributes are optional. They include fields such as geographic coverage, collection granularity, and time range. Adding attributes helps consumers understand more about your data product. Select as many attributes as you like.

Step 4: Details

Provide a short description and a detailed description to help consumers understand what data assets are included in your listing. You can also add resources such as sample notebooks, documentation links, and policy URLs.

Field name	Description
Short description	A short, informative explanation of the dataset that expands on the listing name. This field appears in listing tiles and consumer search results. Recommendations: • Fewer than 100 characters. Cannot exceed 160 characters. • Sentence case (capitalize only the first word and any proper nouns or acronyms).
Description	The detailed description of your data should include a summary of the data and assets being offered in the listing. Basic rich text formatting is supported (that is, bold, <i>italics</i>,  Asr, Genie bullets, and numbered lists), using Markdown syntax.

Field name	Description
	Recommendations: • Include benefits and use cases. • Provide brief guidance about how to use the data and sample use cases. • Include sample datasets and field names. • Specify schemas, tables, and columns. • Use consistent punctuation and syntax. • Add an extra line break between paragraphs. • Don't repeat the attributes that you defined under Add Attribute.
Add notebook	Databricks recommends that you share sample notebooks to demonstrate how best to use the data and to increase listing downloads. Add up to 10 notebooks. For more information about creating notebooks, see Notebook example: Sample notebook and Databricks notebooks. Clicking Add notebook automatically saves the listing.
Documentation	A URL that links to documentation that can help consumers use or understand your dataset (for example, a dataset dictionary).
Privacy policy	A URL that links to your privacy policy. The privacy policy must be publicly accessible and require no login.
License	A URL that links to your license.
Terms of service	A URL that links to your terms of service for the appropriate use of the shared data assets. Terms of service must be publicly accessible and require no login.

MCP server listing fields

 PREVIEW

This feature is in [Public Preview](#).

Model Context Protocol (MCP) servers are tools and APIs that AI agents can use to interact with external systems. You can list MCP servers under the "AI" product category in Marketplace to make them discoverable to AI engineers who are building agents.

When creating an MCP server listing, complete the same basic listing fields as for dataset listings, plus the following MCP-specific fields:

Field name	Description
Authentication type	Specify the authentication method your MCP server requires: • OAuth User to Machine Per User: For OAuth 2.0 authentication flows, requires you to specify a token and authorization endpoint • Bearer Token: For API key or token-based authentication
Host	The hostname or domain where your MCP server is accessible.
Port	The network port number that your MCP server listens on.
Base path (Optional)	The URL path prefix for your MCP server's API endpoints. Leave blank if your MCP server is hosted at the root path.
Tag (Optional)	A version identifier or release tag for your MCP server.
Token endpoint (<i>OAuth User to Machine Per User</i>)	The OAuth 2.0 authorization server endpoint used to obtain access tokens. The authorization code is  Ask Genie

Field name	Description
<i>only)</i>	exchanged for an access token using this endpoint when consumers install this MCP server.
Authorization endpoint (<i>OAuth User to Machine Per User only</i>)	The OAuth 2.0 authorization server endpoint used to get authorization codes. Consumers are redirected to this endpoint to authenticate and authorize access to your MCP server.
OAuth scope (<i>OAuth User to Machine Per User only, optional</i>)	The OAuth 2.0 permissions scope that defines what resources and operations consumers can access.
OAuth credential exchange method (<i>OAuth User to Machine Per User only</i>)	How OAuth 2.0 credentials are exchanged.

Analyze consumer activity using system tables (Public Preview) and dashboards

If you have [system tables](#) enabled in your account, you can use the Marketplace system tables to analyze consumer activity on your listings. For more information, see [Marketplace system tables reference](#).

You can also use the Provider Analytics Dashboard to monitor listing views, requests, and installs. The dashboard pulls data from the system tables. See [Monitor listing usage metrics using dashboards](#).

Resource quotas

Unity Catalog enforces resource quotas on all securable objects, including Marketplace listings per metastore and providers per account. These quotas are listed in [Resource limits](#). If you expect to exceed these resource limits, contact your Databricks account team.

You can monitor your quota usage using the Unity Catalog resource quotas APIs. See [Monitor your usage of Unity Catalog resource quotas](#).

Next steps

- [Manage requests for your data product in Databricks Marketplace](#)
- [Manage existing Databricks Marketplace listings](#)

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Dec 3, 2024**

Manage existing Databricks Marketplace listings

This article describes how to edit, unpublish, delete, and revoke access to Databricks Marketplace listings and listed data products. It is intended for data providers.

Before you begin

You must be a Marketplace admin to manage provider listings. See [Assign the Marketplace admin role](#).

Edit a Marketplace listing

To edit an existing listing:

1. Log into your Databricks workspace.
2. In the sidebar, click the  **Marketplace** icon.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Listings** tab, find the listing you want to modify, click the  kebab menu at the end of the listing row, and select **Edit**.

You cannot change an instantly available listing to one that requires approval, or vice versa.

Unpublish a data product

If you want to remove a listing from the Marketplace UI and disable consumer access requests, you can unpublish it. The listing remains in your workspace in unpublished status, available to be relisted. Any pending requests are deleted.

To unpublish a listing:

1. Log into your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Listings** tab, find the listing you want to unpublish, click the  kebab menu at the end of the listing row, and select **Unpublish listing**.

Delete a data product

If you want to remove a listing from the Marketplace UI, disable consumer access requests, and remove it from your own workspace, you can delete it. You will not be able to relist it. Underlying shares in your workspace are not deleted. Consumers who already have access to the data product in their workspace will continue to have access to the shared data product unless you revoke access to the underlying Delta Sharing share.

To delete a listing:

1. Log into your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Listings** tab, find the listing you want to unpublish, click the  kebab menu at the end of the listing row, and select **Delete**.

Revoke a consumer's access to a data product

To revoke a consumer's access to a data product:

1. Log into your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Consumer requests** tab under **All other requests**, click the request that you want to revoke access to.  Ask Gene

5. Click **Revoke consumer access**.

NOTE

If you revoke a consumer's access to a data product that you have already shared with them, and your listing for that product is still live, the consumer will continue to see the data product as installed on the **My requests** page of the Databricks Marketplace in their workspace, but they will be unable to access the data and unable to request that data product from that listing from that workspace again.

Is this page

as this page helpful?

☐ Yes

☐ No

Last updated on **Apr 6, 2026**

Manage requests for your data product in Databricks Marketplace

This article describes how to manage and monitor consumer requests for your data product in Databricks Marketplace. This article is intended for data providers.

There are two types of Databricks Marketplace listings:

- Listings that are instantly available and free of charge.
- Listings that require your approval because they involve a commercial transaction or require customization (such as a specific geography or demographic sample set, for example).

Manage instantly available listings

You don't need to do anything to let a Databricks Marketplace consumer get access to the data you've shared in your instantly available listings.

To monitor the consumers (recipients, in Delta Sharing terminology) who are using your instantly available listings, you can use the Delta Sharing interfaces. As long as you know the name of the share, you can view all of the recipients who have accessed that share. See [View the recipients who have permissions on a share](#).

You can also use Delta Sharing interfaces to revoke a consumer's access to shared data. See [Revoke recipient access to a share](#).

Manage requests that require your approval

To manage requests that require your approval, use the **Consumer requests** tab in the **Provider console**.

Permissions required: [Marketplace admin role](#). To fulfill a request, you must have the `CREATE RECIPIENT` and `USE RECIPIENT` privileges. See [Unity Catalog privileges reference](#).

To manage a new request:

1. Log into your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. The most recent new requests are listed on the **Overview** tab. To see all requests, go to the **Consumer requests** tab.

New requests are listed before **All other requests**, which includes those that are fulfilled, pending, or denied.

Client type can be **Databricks** (Databricks-to-Databricks sharing) or **Open client** ([Databricks-to-external-platform using a credential file](#)).

5. Click the **Review** button on the request row.
6. Review the requester details.
7. Select an action:
 - **Mark as pending** if your communications with the consumer are still ongoing and any required commercial transactions are incomplete.
 - **Fulfill** if your communication with the consumer and all transactions are complete and you are ready to share the data product. You must **Select a share**, and that share must already be created. See [Create and manage shares for Delta Sharing](#). When you add a share and mark the request fulfilled:

 Ask Genie


- If the consumer is on a Unity Catalog-enabled Databricks workspace, they gain access to that share as a catalog in their Databricks workspace in near-real time.
- If the consumer is not on a Unity Catalog-enabled Databricks workspace, a credential file is generated and made available for them to download. See [Access data products in Databricks Marketplace using external platforms](#).
- **Deny** if you will not share a data product with the consumer. Optionally, you can select a **Reason for denial**. Select **Other** to enter a free-form reason.

8. Click **Update request**.

To manage pending, fulfilled, or denied requests, scroll down to the **All other requests** list. You can view requester details and advance pending requests to another state. You cannot change the status of fulfilled or denied requests.

Monitor listing usage metrics using dashboards

You can use the Provider Analytics Dashboard to monitor consumer usage metrics.



Standard widgets show:

- Listing events (views, requests, and installs) by day over a specified period.
- Seven-day over seven-day trends of views, requests, and installs.
- All listings installed over a specified period, with consumer email address and company.
- Conversions from view to request to install over a specified period
- Conversion rates by listing.

The dashboard pulls data from the Marketplace system tables. You must enable system tables for your workspace in order to use the dashboard. See [Marketplace system tables reference](#).

Requirements

To view and work with the dashboard, you must have:

- The `USE SCHEMA` privilege on the `system.marketplace` schema.
- Access to a compute resource that supports dashboards (compute with standard or dedicated access mode, or a SQL warehouse).

Generate the dashboard in your workspace for the first time

To make the dashboard available in your workspace for the first time, you must open it from the Provider console. In addition to the requirements listed above, you must have the [Marketplace admin role](#).

1. In your Databricks workspace, click  **Marketplace**.
2. On the upper-right corner of the Marketplace page, click **Provider console**.
3. Click the **View Analytics** button.
4. Attach the dashboard to a compute resource that supports dashboards.

For more information about working with dashboards, see [Dashboards](#).

Access the dashboard

Once the dashboard is available in your workspace, you can access it from the Provider console or the Dashboards page:

- To access the dashboard from the **Provider console** (see navigation in the previous section), click the **View Analytics** button.
- To access the dashboard on the **Dashboards** page, go to the sidebar in your workspace and click  **Dashboards**. Search for and select Provider Analytics Dashboard.

Edit the dashboard

You can edit the dashboard as needed. However, Databricks can update the dashboard at any time, and might overwrite your changes. For that reason, Databricks recommends that you make a copy and edit the copy rather than edit the original.

Monitor listing usage metrics using system tables

If you have [system tables](#) enabled in your account, you can use the Marketplace system tables to analyze consumer activity on your listings. For more information, see [Marketplace system tables reference](#).

Is this page

as this page helpful?

☐ Yes

☐ No

Last updated on **Apr 9, 2026**

Create and manage private exchanges in Databricks Marketplace

This article is intended for data providers and describes how to create and manage private exchanges in Databricks Marketplace. Private exchanges enable you to share data products with a defined group of invited consumers. Private exchange listings do not appear in the public marketplace.

What is a private exchange?

A private exchange allows you to make certain data products discoverable only to a specified group of consumers in Databricks Marketplace. Like Delta Sharing, upon which Databricks Marketplace is built, private exchanges give you the ability to share data securely and privately with select recipients. Unlike Delta Sharing, private exchanges have the additional advantage of making data products discoverable by members *before* the data product is shared, so consumers are aware that those data products are available to request. Databricks Marketplace private exchanges also provide consumers with a storefront interface for requesting and accessing data products that might be easier to use than Delta Sharing on its own.

What is a private listing?

Listings are defined as public or private. To share a listing in a private exchange, it must be defined as private. You can add existing private listings to your exchange, and you can create new private listings, assigning them to the exchange when you create the listing. See [Create a listing](#).

You can share free listings that are available instantly or listings that require your approval before the member can access them. You can create new private listings and add them to your exchange at any time. You can also edit a public listing to make it private. See [Create a listing](#).

Before you begin

- By creating a private exchange, you have read and agreed to the [Databricks Marketplace Private Exchange Terms](#).
- To create and manage private exchanges, you need the [Marketplace admin role](#).
- Private exchange members must be Databricks customers and must have access to a workspace that is attached to a Unity Catalog metastore.
- When you add a member to your private exchange, you must enter their metastore's sharing identifier.

The sharing identifier is a string consisting of the metastore's cloud, region, and UUID (the unique identifier for the metastore), in the format `<cloud>:<region>:<uuid>`. For example, `aws:eu-west-1:b0c978c8-3e68-4cdf-94af-d05c120ed1ef`.

To get the sharing identifier, reach out to your contact at the member organization. You might need to tell them how to get the sharing identifier. For instructions, see [Step 1: Request the recipient's sharing identifier](#). You can also point them to step 1 of [Get access in the Databricks-to-Databricks model](#) for instructions.

Set up a private exchange

You create the private exchange using the provider console in Databricks Marketplace. You can add members and listings when you create the exchange or after.

Permission required: [Marketplace admin role](#).

To create a private exchange:

1. Log into your Databricks workspace.

2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Exchanges** tab in the provider console, click **Create exchange**.
5. On the **Create exchange** dialog, enter the name of your exchange and click **Create**.

Use a human-friendly name that helps you and other users who are managing exchanges recognize the purpose of the exchange.

6. On the **Exchanges** tab, find the exchange and click the name.
7. Add the members you want to share with.

- i. On the **Members** tab, click **Add member**.
- ii. On the dialog, add a human-readable name for the member and enter the **Sharing identifier**.

To learn how to get the sharing identifier from the member organization, see [Before you begin](#).

- iii. Click **Add member**.
8. Add the listings you want to share.
 - i. On the **Listings** tab, click **Add listing**.
 - ii. On the dialog, select one of the existing private listings from the drop-down list and click **Add listing**.

You can also create a new private listing or edit a public listing to make it private.

- i. Click **create a new listing** to go to the **New listing** page.
- ii. When you create the listing, select the **Private exchange** option and select the exchange from the drop-down list. See [Create a listing](#).

Edit members or remove them from a private exchange

Permission required: [Marketplace admin role](#)

To edit or remove a member:

1. Log into your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Exchanges** tab in the provider console, find and click the exchange name.
5. On the **Members** tab, find the member, and click the  kebab menu at the end of the member row.

To edit the member, select **Edit**. You can update both the member name and sharing identifier.

To remove the member, select **Remove** on the  kebab menu.

Remove listings from a private exchange

Permission required: [Marketplace admin role](#).

To remove a listing from an exchange:

1. Log into your Databricks workspace.
2. In the sidebar, click  **Marketplace**.
3. On the upper-right corner of the Marketplace page, click **Provider console**.
4. On the **Exchanges** tab in the provider console, find and click the exchange name.
5. On the **Listings** tab, find the listing, click the  kebab menu at the far right, and select **Remove**.
6. On the confirmation dialog, click **Remove**.

 Ask Genie

You can also remove the link between a listing and an exchange by editing the listing and removing the exchange from the field under **Private exchange**.

Manage member requests for data products

You manage member requests for data products in private exchanges the same way that you do for data products in the public marketplace. See [Manage requests for your data product in Databricks Marketplace](#).

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Oct 10, 2023**

Databricks Marketplace provider policies

Thank you for your interest in becoming a Provider on the Databricks Marketplace (“**Marketplace**”). This Marketplace Policy (“**Policy**”) applies to all Providers and their Listing(s) of Product(s) in the Marketplace. This Policy is part of your Databricks Data Partner Program Agreement (“**Agreement**”), as well as any other terms that reference this Policy. Capitalized terms not defined in this Policy have the meanings given elsewhere in the Agreement (or the terms to which these apply). For convenient reference: Your Agreement consists of this Policy, the Marketplace Documentation, the Databricks Data Partner Program Overview (“**Program Overview**”), and the Terms, all as more fully defined in the Program Overview.

Providers must comply at all times with the policies described and/or referenced in this document in order to participate in the Marketplace. If a Provider fails to meet any of the enumerated policies at any given time, Databricks may, among other things, remove the Provider and/or Listing(s) from the Marketplace. We may update this Policy from time to time as set forth in the Agreement.

General Provider requirements

- To publish a Listing for Products in the Marketplace, you must provide all information required by Databricks. You must:
 - Clearly and accurately describe who you are. In other words, you may not impersonate or imply that a relationship or affiliation that does not exist; and
 - Only use trademarks for which you own or are authorized to use.
- You must have all necessary rights to sell or share the Product(s) you intend to offer.
- Only one profile is allowed per Provider, unless otherwise approved in writing by Databricks.

- Your profile must be accurate and kept up to date.
- Manipulation of your position in the Marketplace is expressly forbidden.
- All Providers are expected to promptly respond to all Consumers as follows:
 - Within three business days to inquiries from prospective Consumers requesting access to your Product; and
 - All other inquiries must be responded to within five business days.
- Providers must deliver the Product as advertised in the Listing and provide Consumers notice when a change is made to the Listing that may affect them.

Listing and Product requirements

Listings

- Any Listing published in the Marketplace must provide all information required by Databricks including but not limited to:
 - An accurate description of what is being offered including a brief and clear explanation of the Product(s);
 - Your Provider/Consumer Terms of use, such as those describing the license grant, any other pertinent terms and conditions, and restrictions (if any) applicable to the Product being offered;
 - Accurate documentation;
 - Update frequency;
 - Whether a Product consists of or includes Data Assets (for example, notebooks or dashboards);
 - A description of Personal Data contained in the Product, if applicable; and
 - Keeping all links up-to-date.
- Listings may not:
 - Include Personal Data;
 - Include, advertise, or promote illegal content;
 - Encourage or imply illegal, threatening, or violent uses; or
 - Advertise or promote anything other than the Product(s) in the Listing.
- Products and Listings may not:
 - Infringe any third party intellectual property rights;
 - Contain or provide access to inappropriate, harmful, or offensive content; or

- Encourage Consumers to access or use the Product outside of the service or functionality set forth in the Listing unless otherwise agreed upon by Databricks.
- Any material changes to your Listing may be reviewed by Databricks, and Databricks reserves the right to remove the Listing.

Products

- For clarity – as stated in the Program Overview, if your Product includes Personal Data, you must be legally authorized to share such Personal Data. However, a Product must never include Sensitive Personal Data.
- If you are offering a Dataset, you must update the Dataset at the most frequent logical refresh rate. Ideally, this is as close to real time as possible.
- If your Product shares anonymized or aggregated data, you must structure the data in such a way as to ensure that it remains anonymous even when combined with other data.
- In addition to the requirements in this Policy, your Product(s) must also comply with the descriptions and commitments made in your Provider/Consumer Terms.
- Any notebooks accompanying queries and/or code content in your materials must work for all Consumers and produce the results advertised in your materials.

Data shared by Databricks with Providers

- Any usage data we share with you about your Listing is confidential and only for your internal use. You may not use this data for benchmarking.
- Any Personal Data we share with you may only be used in connection with marketing and selling your Product(s) in your Listing(s) in a manner consistent with your agreement with Databricks, your obligations to Consumers, and in accordance with your privacy policy as well as any other laws that may be applicable.

Delisting and discontinuing a Product

- You may delist a Listing (cease to offer) at any time, such that new Consumers cannot discover the Listing; and
- You may cease supplying a Product so long as you:

- Comply with the notice requirements set forth in the Provider/Consumer Terms but in no event shall said notice be less than 30 days; and
- Allow any current Consumer(s) continued access to the Product for the remainder of the current term you granted under your Provider/Consumer Terms with the Consumer(s).

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Apr 21, 2026**

What is Databricks Clean Rooms?

This page introduces Clean Rooms, a Databricks feature that uses Delta Sharing and serverless compute to provide a secure and privacy-protecting environment where multiple parties can work together on sensitive enterprise data without direct access to each other's data.

Requirements

To be eligible to use clean rooms, you must have:

- A workspace that is enabled for serverless compute. See [Serverless compute requirements](#).
- A workspace that is enabled for Unity Catalog. See [Enable a workspace for Unity Catalog](#).
- Delta Sharing enabled for your Unity Catalog metastore. See [Enable Delta Sharing on a metastore](#).

How does Clean Rooms work?

When you create a clean room, you create the following:


- A securable clean room object in your Unity Catalog metastore.
- The “central” clean room, which is an isolated ephemeral environment managed by Databricks.
- A securable clean room object in your collaborator's Unity Catalog metastore.

Tables, volumes (non-tabular data), views, and notebooks that either co-located or are shared with the central clean room only, using Delta Sharing.

 Ask Genie

Sharing.

Collaborators cannot see the data in other collaborators' tables, views, or volumes, but they can see column names and column types, and they can run approved notebook code that operates over the data assets. The notebook code runs in the central clean room. Notebooks can also generate [output tables](#) that let your collaborator temporarily save read-only output to their Unity Catalog metastore so they can work with it in their workspaces.



How does Clean Rooms ensure a no-trust environment?

The Databricks Clean Rooms model is “no-trust.” All collaborators in a no-trust clean room have equal privileges, including the creator of the clean room. Clean Rooms is designed to prevent the running of unauthorized code and the unauthorized sharing of data. This trust is enforced explicitly by requiring all collaborators to approve a notebook before it can be run. While you can upload a notebook to yourself, it must be approved by the other collaborators before running. See [Approve a notebook in a clean room](#).

Additional safeguards or restrictions

The following safeguards are in place in addition to the explicit notebook approval process mentioned above:

- After a clean room is created, it is locked down to prevent new collaborators from joining the clean room.
- If any collaborator deletes the clean room, the central clean room is void and no clean room tasks can be run by any user.
- Each clean room is limited to ten collaborators.
- You cannot rename the clean room.

The clean room name must be unique in every collaborator's metastore, so that all collaborators can refer to the same clean room unambiguously.

- When a collaborator adds comments to a clean room securable in their workspace, these comments are not propagated to other collaborators.

What is shared with other collaborators?

- Clean room name.
- Cloud and region of the central clean room.
- Your organization name (which can be any name you choose).
- Your clean room sharing identifier (global metastore ID + workspace ID + user email address).
- Aliases of shared tables, views, or volumes.
- Column metadata (column name or alias and type).
- Notebooks (read-only).
- [Output tables](#) (read-only, temporary).
- Clean room events system table.
- Run history, including:
 - The name of the notebook that is being run
 - Collaborator that ran the notebook (not user).
 - The state of the notebook run.
 - The start time of the notebook run.

What is shared with the central clean room?

- Everything that is listed in the previous section.
- Read-only tables, volumes, views, and notebooks.

Tables, views, and volumes are registered in the central clean room's metastore with any supplied aliases. Data assets are shared throughout the lifecycle of the clean room.

Clean Rooms FAQ

The following are frequently asked questions about clean rooms.

How is my data managed in a clean room?

The central clean room is managed by Databricks. In the central clean room:

- Neither party has admin privileges.
- Only metadata is visible to all parties.
- Each party can add data to the central clean room.
- Clean rooms use Delta Sharing to share data securely to the clean room, but not between participants. See [What is Delta Sharing?](#).

How is my data kept private?

Central clean rooms run in an isolated, Databricks-managed serverless compute plane hosted in a cloud provider region that the clean room creator chooses.

Clean rooms provide:

- **Code approval:** The clean room creator and collaborators can share tables and volumes with the central clean room but can only run notebooks uploaded by the other party. You can review the code added by the other party before approving. If you run a notebook added by another party, you implicitly approve the code.
- **Version control:** Clean rooms notebooks have version control to ensure that all parties can only run fully-approved notebooks. Only the most recent version of a

notebook can be run. You can use the clean rooms system table to see which version of the notebook was run and monitor any changes made.

- **Restricted access:** When you create a clean room, you can use serverless egress control to manage outbound network connections. If you restrict access from your clean room, access to unauthorized storage is blocked. See [What is serverless egress control?](#).

To learn more about security and the serverless compute plane, see [Serverless compute plane networking](#).

How are actions recorded?

Clean room actions taken by you or your collaborators are recorded in the clean room events system table. These records include detailed metadata about the specific action taken. See [Clean room events system table reference](#).

Clean room actions are also recorded in your account's audit log under the service `clean-room`. See [Clean Rooms events](#).

When a collaborator modifies permissions on a clean room securable within their metastore, these changes are recorded in their audit logs under the `unityCatalog` service.

How does billing work?

To learn more about Databricks Clean Rooms pricing, see [link](#).

Limitations

The following limitations apply:

- No service credential Scala libraries included in the required Databricks Runtime version.
- If you use default storage tables in a clean room, Databricks recommends enabling the **Delta Sharing for Default Storage – Expanded Access** feature ([Beta](#)) in your account console.

Resource quotas

Databricks enforces resource quotas on all Clean Room securable objects. These quotas are listed in [Resource limits](#). If you expect to exceed these resource limits, contact your Databricks account team.

You can monitor your quota usage using the Unity Catalog resource quotas APIs. See [Monitor your usage of Unity Catalog resource quotas](#).

Get started

- [Create clean rooms](#)
- [Run notebooks in clean rooms](#)
- [Manage clean rooms](#)
- [Work with Databricks clean rooms as an invited collaborator](#)

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Apr 21, 2026**

Create clean rooms

This page describes how to create a Databricks clean room using the UI. A clean room is a secure environment for collaborative data analysis.

Key features and limitations:

- **Secure collaboration:** Clean rooms enable multiple parties to work together on sensitive enterprise data without direct access to each other's raw data.
- **Collaborator capacity:** A clean room can have ten parties, including the creator and up to nine other collaborators.
- **Metastore limit:** Databricks enforces resource quotas on all clean room securable objects. See [Resource limits](#).

To use the REST API, see [Create a clean room](#).

Before you begin

The privileges needed to use clean rooms vary depending on the task:

Task	Required privileges
View a clean room	Must be the owner of the clean room, or have one of the following privileges on the clean room: <code>MANAGE</code> , <code>MODIFY CLEAN ROOM</code> , <code>EXECUTE CLEAN ROOM TASK</code> , or <code>BROWSE</code> .
Update the owner of a clean room	Must be the owner of the clean room, or have the <code>MANAGE</code> privilege on the clean room.
Add or remove data assets in	Must be the owner of the clean room or have the <code>MODIFY CLEAN ROOM</code> privilege on the clean room. If you are not the c

Task	Required privileges
a clean room	clean room, you and the clean room owner must have <code>SELECT</code> on any table or view and <code>READ VOLUME</code> on any volume that you add, along with <code>USE CATALOG</code> and <code>USE SCHEMA</code> on the parent catalog and schema.
Add or remove notebooks in a clean room	For the notebook's uploader: • If they are the designated runner of the notebook, they must have the <code>EXECUTE CLEAN ROOM TASK</code> or <code>MODIFY CLEAN ROOM</code> privileges. • If the notebook is runnable by a collaborator, the uploader must have the <code>MODIFY CLEAN ROOM</code> privilege.
Update a comment in a clean room	Must be the owner of the clean room or have the <code>MODIFY CLEAN ROOM</code> privilege on the clean room.
Grant access to a clean room	Must be the owner, or have the <code>MANAGE</code> privilege on the clean room.
Delete a clean room	Must be the owner or have the <code>MANAGE</code> privilege on the clean room.

Beyond the task-specific privileges, when a clean room is shared, the email address attached to the sharing identifier automatically becomes the owner in your collaborator's organization. See [Step 1. Request the collaborator's sharing identifier](#).

To learn about permission requirements for updating clean rooms and running tasks (notebooks) in clean rooms, see [Manage clean rooms](#) and [Run notebooks in clean rooms](#).

NOTE

The central clean room can have a maximum of two other non-AWS regions among its collaborators.

 Ask Genie

Step 1. Request the collaborator's sharing identifier

Before you can create a clean room, you must have the Clean Room sharing identifier of the organizations you will be collaborating with. The sharing identifier is a string that consists of the organization's global metastore ID + workspace ID + the contact's username (email address). Your collaborators can be in any cloud or region.

Reach out to your collaborators to request their sharing identifier. They can get the sharing identifier using the instructions in [Find your sharing identifier](#).

Step 2. Create a clean room

To create a clean room, you must use Catalog Explorer.

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.
3. Click **Create Clean Room**.
4. On the **Create Clean Room** page, enter a user-friendly name for the clean room.

The name cannot use spaces, periods, or forward slashes (/).

Once it's saved, the clean room name cannot be changed. Use a name that potential collaborators will find useful and descriptive.

5. Select the cloud provider and region where the central clean room will be created.

The cloud provider must match your current workspace, but the region can be different. Consider your organization's data residency or other policies when you make your selection.

6. Each clean room can have up to ten collaborators. Enter the **Clean Room sharing identifier** for each collaborator. See [Step 1. Request the collaborator's sharing identifier](#).

Create clean room

[Send feedback](#)

General

Set up the clean room

Clean Room name*

Region*

This will be the cloud and region for your secure central clean room.

us-west-2

Collaborator details

Collaborators cannot be modified after the clean room is created.

Clean Room sharing identifier* ⓘ

aws-eu-west-1:abcde

+ Add collaborator

You can test your clean room before full deployment by using either your sharing identifier or the identifier of another user in your current metastore. Doing so creates two clean rooms in your current metastore. For example, if you create a clean room titled `test_clean_room`, a second clean room named `test_clean_room_collaborator` also appears. Running notebooks with a collaborator in the same metastore functions the same as with an external collaborator. See [Run notebooks in clean rooms](#).

7. Make a note of the catalog names assigned to you and your collaborators.

All data assets added to the clean room will appear under that catalog in the central clean room, and can be referenced using that catalog in the Unity Catalog three-level namespace (`<catalog>.<schema>.<table-etc>`).

8. Select the network access policy type. This cannot be changed after the clean room is created.

- **Full Access:** Unrestricted outbound internet access.
- **Restricted Access:** This limits outbound access to internet destinations that you specify. See [Network policy overview](#) and [Manage network policies for serverless egress control](#).

Security

Security settings cannot be modified after the clean room is created.

Network access

Control which network destinations can be accessed from your clean room. Access to the cloud storage destinations backing your clean room assets will be automatically allowed.

☐ Full internet access

All internet destinations will be accessible from your clean room.

☒ Restricted internet access

Only internet destinations added to this policy will be accessible from your clean room.

+ Add destination

Restricted access can delay asset availability for up to ten minutes.

After you create the clean room, you can view the network access policy in the **Security** tab.

1. Click **Create Clean Room**.

If your current workspace is set to the HIPAA compliance security profile, then when you create a clean room, that setting is applied to the central clean room.

Collaborators must access the clean room from a workspace with the same security profile. See [Compliance security profile](#).

When creating a clean room, its security profile is inherited from your current workspace settings. This applies if your workspace is configured for HIPAA or FedRAMP Moderate compliance. FedRAMP Moderate is supported only when the clean room region is set to `us-west-2` or `us-east-1`. All collaborators must access the clean room from a workspace configured with the identical security profile to ensure consistent compliance.

Step 3. Add data assets and notebooks to the clean room

Both the creator and the collaborators can add tables, volumes, views, and notebooks to the clean room.

NOTE

The following instructions assume you are returning to an already-created clean room to add assets. If you just created a clean room for the first time, a wizard walks you through adding data assets and notebooks. The actual UI for adding these assets is the same, regardless of whether you are guided by the wizard or not.

To add notebooks:

1. Click the **+ Add notebooks** button and browse for the notebook you want to add.

 Ask Genie

2. Name the notebook.

3. Select which collaborator can run the notebook. Select **You** to run the notebook yourself.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Add notebook

Notebooks will be uploaded as static files to the clean room.

Notebook
/Users/test.user@collaborator1.com/Analysis Notebook

Notebook name
Analysis Notebook Preview

Notebook runner
Collaborator 1 (You) ▼

- ✓ Collaborator 1 (You)
- Collaborator 2
- Collaborator 3
- Collaborator 4

You can optionally give the notebook an alternative **Notebook name**.

Notebooks that you share in clean rooms query data and run data analysis workloads on the tables, views, and volumes that you and any other collaborators have added to the clean room.

If you share a notebook that includes results, those results are shared with your collaborators.

You can use a notebook to create *output tables* that are temporarily shared to your collaborator's metastore when they run the notebook. See [Create and work with output tables in Databricks Clean Rooms](#).

To use a test dataset, download our [sample notebook](#).

❗ IMPORTANT

Any notebook references to tables, views, or volumes that were added to the clean room must use the catalog name assigned when the clean room was created ("creator" for data assets added by the clean room creator, and "collaborator" for data assets added by the invited collaborator). For

example, a table added by the creator could be named `creator.sales.california`.

Likewise, verify that the notebook uses any aliases assigned to data assets that were added to the clean room.

To add assets:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms >** button.
3. Find and click the name of the clean room you want to update.
4. Click **+ Add data assets** to add tables, volumes, or views.
5. Select the data assets you want to share and click **Add data assets**.

When you share a table, volume, or view, you can optionally add an alias. The alias name is the only name visible in the clean room.

When you share a table, you can optionally add partition clauses that enable you to share only part of the table. For details about how to use partitions to limit what you share, see [Specify table partitions to share](#).

 NOTE

If you use default storage tables in a clean room, Databricks recommends enabling the **Delta Sharing for Default Storage – Expanded Access** feature ([Beta](#)) in your account console.

 NOTE

Federated table sharing is in [Beta](#). See [Add foreign schemas or tables to a share](#).

[Data sharing](#)[Clean Rooms](#)

Run notebooks in clean rooms

Last updated on **Apr 30, 2026**

Run notebooks in clean rooms

This page describes how to run notebooks in clean rooms. Notebooks are the interface that collaborators use to run data analysis in collaboration.

To learn how to add a notebook to a clean room, see [Create clean rooms](#).

Before you begin

To manage and run a notebook in a clean room, you must:

[Ask Genie](#)

Task	Requirement	Required privileges	Description
Managing a notebook	No unique requirements	For the notebook's uploader: - If they are the designated runner, they must have <code>EXECUTE CLEAN ROOM TASK</code> or <code>MODIFY CLEAN ROOM</code> privileges. - If a collaborator is the designated runner, the uploader must have the <code>MODIFY CLEAN ROOM</code> privilege.	General management tasks for a notebook in the clean room.
Running a notebook	Every collaborator, except the uploader, must approve the notebook. A designated runner runs the notebook.	<code>EXECUTE CLEAN ROOM TASK</code>	If the runner of the notebook did not upload it, they must approve the notebook before it can be run. This explicit approval can be automated through the default auto-approval rule. See Auto-approval rules .
Approving or rejecting a	None	<code>MODIFY CLEAN ROOM</code>	Allows you to  Ask Genie approve or reject

Task	Requirement	Required privileges	Description
notebook			notebooks before they can be run.
Managing auto-approval rules	Auto-approvals can only be applied to notebooks authored by collaborators other than the runner.	Clean room owner or <code>MANAGE</code> privilege.	Controls the automatic approval of notebooks.
Collaborator capacity	A clean room can include up to 10 collaborators.	None	This includes the creator and up to 9 other collaborators.

The creator is automatically assigned as the owner of the clean room in their Databricks account. The collaborator organization's metastore admin is automatically assigned ownership of the clean room in their Databricks account. You can transfer ownership. See [Manage object ownership](#).

Uploading a notebook and designating runners

When a collaborator adds a notebook to the clean room:

- They are considered the uploader of that notebook in the clean room.
- When adding a notebook, the uploader must designate which collaborator is the designated runner of the notebook. The designated runner is permitted to run the notebook.
- There can only be one designated runner.
- You can assign yourself as the designated runner of the notebook.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Add notebook

Notebooks will be uploaded as static files to the clean room.

Notebook

/Users/test.user@collaborator1.com/Analysis Notebook 

Notebook name

Analysis Notebook Preview

Notebook runner

Collaborator 1 (You) ▼

- ✓ Collaborator 1 (You)
- Collaborator 2
- Collaborator 3
- Collaborator 4



See [Step 3. Add data assets and notebooks to the clean room](#)

Approve a notebook in a clean room

Every notebook requires approval from all collaborators, except the uploader, before executing it.

To check the review status of a notebook:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms >** button.
3. Select the clean room from the list.
4. Select the notebook you want to run.
5. Click the **People** icon on the right to expand the collaborator section of the notebook details page.
6. The **Reviews** section provides a clear overview of which collaborators have approved, rejected, or have yet to review the notebook.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Analysis Notebook

Preview Runs

```
%sql
select * from collaborator_1.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_2.ads_demo.audience_data;
```

```
%sql
select * from collaborator_3.audience_overlap.advertiser_customers;
```

Collaborators New

Permissions

Owner	Collaborator 2
Runner	Collaborator 4

Reviews 2/4 approved

- ⌵ **Collaborator 1 (You)**
Pending review
- ✅ **Collaborator 2**
Added on May 12, 2025, 03:55 PM
System-generated
- ⌵ **Collaborator 3**
Pending review
- ✅ **Collaborator 4**
Auto-approved on May 12, 2025, 03:55 PM

A difference view is available when a notebook has been changed from a previous version.

When running a notebook version, the following is true:

- You can only run the latest version of a notebook.
- You can only approve or reject the latest version of a notebook.
- Modifying a notebook by adding new content or changing the runner designation creates a new version, which resets the review state for all collaborators.

To access the difference view for a notebook:

1. Click the **Clock icon** on the right to expand the version history of the notebook details page.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Analysis notebook 3

Preview Runs

You are viewing the latest version from Apr 18, 2025, 10:29 AM. [Exit](#)

```
%sql
select * from collaborator_2.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_1.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_1.ads_demo.audience_data;
```

```
%sql
select * from collaborator_2.ads_demo.audience_data;
```

```
%sql
select * from collaborator_3.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_3.audience_overlap.advertiser_customers;
```


Version history

- Apr 18, 2025, 10:29 AM
✅ You approved
- Apr 18, 2025, 09:58 AM
✅ You approved
- Apr 18, 2025, 07:11 AM
✅ You approved
- Apr 16, 2025, 03:43 PM
✅ You approved

Ask Genie

To approve or reject a notebook:

1. Click the **Review** button at the upper right.
2. Choose to **Approve** or **Reject** the notebook.



The screenshot shows a modal dialog box titled "Update review for 'Analysis Notebook'" with a close button (X) in the top right corner. The dialog is divided into three sections:

- 1. Review notebook content**: A text area with the instruction "Confirm that the use of any private libraries and notebook parameters are correct."
- 2. Review notebook runner**: A dropdown menu showing "Collaborator 4" with a document icon.
- 3. Add a comment (optional)**: A large text input field.

At the bottom of the dialog are two buttons: "Reject" with a red 'X' icon and "Approve" with a green checkmark icon.

You can update your review up to nine times, for a maximum of ten reviews per collaborator on each notebook version. However, you can always reject a notebook, even if the review limit has been reached.

Auto-approval rules


Collaborators can set up auto-approval rules for their clean room. With auto-approvals, the following is true:

- You can create auto-approval rules only for notebooks uploaded by other users, not your own. Self-authored notebooks do not require your own approval if you are the designated runner.
- In 2-person clean rooms, you can auto-approve notebooks authored by the other collaborator.
- In clean rooms with more than two collaborators, you can auto-approve notebooks authored by **Anyone** or a specific collaborator.

- Each auto-approval rule designates a single runner for the approved notebook.
- Auto-approval is the default for notebooks uploaded by another collaborator when you are the designated runner.
- Auto-approvals are optional for clean rooms where the designated runner is another collaborator.
- You can add, update, or remove auto-approvals at any time.

To manage auto-approvals:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.
3. On the **Clean Rooms** page click the **Gear icon**
4. Select **Configure auto-approvals**.



Catalog Explorer > Clean Rooms >

 **multi-collaborator-clean-room**

Overview Runs Permissions Security


Filter assets Managed by: All collaborators

Notebooks (query data and run analysis)

Name	Status	Managed by	Added on	
 Analysis notebook 1	2/4 approved	e2-dogfood-clean-room-us-west-2	Apr 23, 2025, 10:20 AM	 Pending
 Analysis notebook 2	Rejected	e2-dogfood-clean-room-us-west-2	May 01, 2025, 11:07 AM	 Rejected

Configure auto-approvals

5. Configure automatic notebook approval based on matching permission rules.



Run a notebook in a clean room

To run a notebook, check all necessary approvals are in place and that you are the designated runner. Then:

1. Under **Notebooks**, click the notebook to open it in preview mode.
2. Click the **Run** button.
 - You can only run notebooks for which you are the designated runner and which have been approved.
 - It is not necessary to manually approve notebooks that you authored yourself to run them.
 - You can reject a notebook you authored. You must re-approve it before it can be run.
 - You can manage auto-approvals for each notebook. See [Auto-approval rules](#).
3. (Optional) On the **Run notebook with parameters** dialog, click **+ Add** to pass parameter values to the notebook job task.
4. Review the notebook.
5. Click **Run**.
6. Click **See details** to view the progress of the run. Alternatively, you can view run progress by going to **Runs** on this page or by clicking **Jobs & Pipelines** in the workspace sidebar and going to the **Job runs** tab.

7. View the results of the notebook run. The notebook results appear after the run is complete. To view past runs, go to **Runs** and click the link in the **Start time** column.

Even when all collaborators approve a notebook, only the collaborator designated by the uploader as the runner can run it.

Notebook parameters

The following parameters are automatically passed into the clean room notebook at runtime:

- `cr_central_id`: The central clean room ID.
- `cr_runner_global_metastore_id`: The global metastore ID of the designated runner.
- `cr_runner_alias`: The collaborator alias of the designated runner.
- `cr_<alias>_input_catalog`: The catalog that stores data shared by a specific collaborator, where `alias` represents their clean room alias. For two-party clean rooms created in the UI, this defaults to creator or collaborator, but it is customizable via the API. You can use `cr_<alias>_input_catalog` parameter for local tests with sample tables.
- `cr_output_catalog`: Defines the catalog where you create new output tables. Used in combination with `cr_output_schema`. See [Create an output table](#).
- `cr_output_schema`: Defines the schema where you create output tables. Used in combination with `cr_output_catalog`. See [Create an output table](#).

The above notebook parameters are automatically available as widget values during execution.

- **Python cells**: Reference them using `dbutils.widgets.get`. For example: `dbutils.widgets.get("cr_central_id")`.
- **SQL cells**: Access them with `select :name`. For example: `select :cr_central_id`.

Share notebook output using output tables

Output tables are temporary read-only tables generated by a notebook run and shared to the notebook runner's metastore. If the notebook creates an output table, the notebook runner can access it in an *output catalog* and share it with other users in their workspace. See [Create and work with output tables in Databricks Clean Rooms](#).

Use Lakeflow Jobs to run clean room notebooks

You can build complex, recurring workflows around your clean room assets using Databricks workflows. For example, you can create a recurring workflow that runs a clean room notebook task, then runs a task that immediately updates a report based on the clean room's output.

The following features facilitate complex clean room workflows:

- **Clean Room notebook task type:** Directly select and run a clean room notebook as a dedicated workflows task. See [Run notebooks in clean rooms](#).
- **Databricks Supplied output values:** All clean room notebooks tasks make the `{{tasks.<your_task_name>.output.catalog_name}}` and `{{tasks.<your_task_name>.output.schema_name}}` dynamic value references available to all downstream tasks. These values can be used for setting up workflows where a task is automatically passed the path to an upstream clean room notebook tasks's output schema. See [Supported value references](#)
 - `output.catalog_name` is automatically populated with the clean room's output catalog name.
 - `output.schema_name` is automatically populated with the clean room notebook task's dynamically-generated output schema name.
- **Lakeflow Jobs:** Use **Task values** that pass job parameter values to clean room notebooks or capture clean room notebook output and pass that output to other workflow tasks. See [Task values](#) to pass information between tasks.
- **Task values:** Like regular notebook tasks, clean room notebook tasks can set **Task values** that are passed to downstream workflow tasks. Use the syntax: `dbutils.jobs.taskValues.set(key="key", value="value")`. See more about [Task values](#).

Is this page

Was this page helpful?

☐ Yes	☐ No
Send feedback	

Last updated on **Dec 18, 2025**

Create and work with output tables in Databricks Clean Rooms

This page introduces output tables, which are temporary read-only tables generated by a notebook run and shared to the notebook runner's Unity Catalog metastore. This article describes how to use a notebook to create output tables and how runners can read these output tables in their Unity Catalog metastore.

Overview of output tables

Output tables let you temporarily save the output of notebooks that are run in a clean room to an output catalog in your Unity Catalog metastore, where you can make the data available to members of your team who don't have the ability to run the notebooks themselves. You can also use Lakeflow Jobs to run notebooks and perform tasks on output tables. Combined with the Clean Room notebook task type and support for task values, output tables let you create complex workflows that depend on Clean room notebooks.

Output tables are read-only.

Only the specific principal (user, group, or service principal) who runs the notebook has default read access to the output table. There is no write access. A metastore admin can grant read access to other principals in their Databricks account, using standard Unity Catalog privileges.

Output tables are stored for 30 days in the central clean room's default storage location and shared to the runner's metastore using Delta Sharing. If you want to keep an output table for more than 30 days, you must copy it to local storage.

Each notebook run creates a new schema in the output catalog. New runs cannot append an existing output table.

NOTE

Collaborators in Databricks on all three clouds—AWS, Azure, and Google Cloud—can share notebooks that create output tables and can read output tables that are generated when they run shared notebooks.

Create an output table

To create an output table, use the parameters `cr_output_catalog` and `cr_output_schema` in the three-part table namespace. Each run of the notebook produces a new schema.

In the following example, the notebook cell creates an output table called `overlapping_users` in the runner's output catalog that lists the users whose email address shows up in both the `collaborator.advertiser.profiles` and `creator.publisher.profiles` tables.

SQL

```
CREATE TABLE identifier(:cr_output_catalog || '.' || :cr_output_schema ||  
'overlapping_users') AS  
SELECT collab_profiles.*  
FROM collaborator.advertiser.profiles AS collab_profiles  
JOIN creator.publisher.profiles AS creator_profiles  
ON collab_profiles.email = creator_profiles.email
```

Read an output table

Output tables appear in a shared catalog in the notebook runner's metastore. In the Catalog Explorer **Catalog** pane, they appear in the **Shared** catalogs list.

Reading an output table is like reading any other table in Unity Catalog. You must have `SELECT` on the table, `USE CATALOG` on the shared output catalog, and `USE SCHEMA` on  Ask Genie

the automatically-generated schema. The user who ran the notebook that created the table has these permissions by default.

NOTE

Deleting a clean room removes all output tables and historical data from the output catalog.

Before you begin

This section describes cloud, configuration, and compute requirements for reading output tables.

Shared output catalog requirement

Before you can read output tables, a user must create the catalog that holds them. You only need to do this once per clean room. The owner of the clean room has permission to read and manage the output catalog by default.

Permissions required: EXECUTE CLEAN ROOM TASK

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms >** button.
3. Select the clean room from the list.
4. In the right pane, under **Output**, click **Create catalog**.
5. Enter an **Output catalog name** or accept the default, which is `<clean-room-name>_output`.

The output catalog appears in the list of **Shared** catalogs in your Catalog Explorer **Catalog** pane. Each clean room that you participate in can have one shared output catalog in your metastore.

Compute requirements

Queries on output tables require serverless compute. See [Connect to serverless compute](#).

Permissions required to read an output table

The user who ran the notebook that created the output table and the owner of the clean room have permission to read and manage the output table by default. All other users must have the following permissions granted to them:

- `SELECT` on the table
- `USE CATALOG` on the output catalog
- `USE SCHEMA` on the output schema

Run the notebook

To generate shared output tables in your output catalog, a user with access to the clean room must run the notebook. See [Run notebooks in clean rooms](#). Each notebook run creates a new output schema and table.

TIP

You can use Lakeflow Jobs to run notebooks and perform tasks on output tables, enabling complex workflows. See [Use Lakeflow Jobs to run clean room notebooks](#).

Find and view an output table

The user who runs the notebook that creates the output table can find a link to the output table on the notebook run history and run details pages in the **Clean Rooms** UI. In both cases, the link is in the **Output schema** field. See [Monitor clean room notebook runs](#).

Run history:

Catalog Explorer > Clean Rooms >						
 test-clean-room						
Overview Runs Permissions						
Notebook ▾ Run status ▾ Run by ▾ 1 job run						
Status	Start time	Run by	Notebook	Duration	Output schema	Expiration
	Sep 18, 2024, 03:08 PM	acme_co (You)	Output Tables - SQL	2m 38s	output_schema_57316933	Expires in 21 days

Run details:

Workflows > Runs > Run 251799029024345 >

CleanRoomsNotebookTask run Succeeded

Edit task Repair

Output

✓ 46.24 seconds 1: Create Output Table

```
%sql
CREATE TABLE identifier(
  :cr_output_catalog ||
  '.' ||
  :cr_output_schema ||
  '.overlapping_users'
) AS
SELECT collab_profiles.*
FROM collaborator.advertiser.profiles AS collab_profiles
JOIN creator.publisher.profiles AS creator_profiles
ON collab_profiles.email = creator_profiles.email;
```

_sqlidf: pyspark.sql.connect.dataframe.DataFrame = [num_affected_rows: long, num_inserted_rows: long]

Table

num_affected_rows num_inserted_rows

No rows returned

0 rows | 46.24 seconds runtime

This result is stored as _sqlidf and can be used in other Python and SQL cells.

Task run

Details Metrics

Details

Job ID 672949637986455

Job run ID 349315244860569

Task run ID 573169339677688

Run as

Started 09/18/2024, 03:08:01 PM

Ended 09/18/2024, 03:10:44 PM

Duration 2m 43s

Queue duration -

Status Succeeded

Lineage No lineage information for this job. [Learn more](#)

View run events View run libraries

Clean Room

Notebook test-clean-room/Output Tables - SQL

Output schema test-clean-room_output/output_schema_...

Expiration Expires in 21 days

You can also find the output catalog in the list of **Shared** catalogs in your Catalog Explorer **Catalog** pane.

Limitations

In addition to the requirements listed in [Overview of output tables](#) and [Before you begin](#), output tables have the following limitations:

- Output tables are only supported when the clean room was created after the output table feature was released.
- Only tables are supported. Volumes and views, for example, are not.
- You are limited in the number of output tables each notebook can support. See [Resource limits](#).

Is this page
as this page helpful?

👍 Yes

👎 No

Ask Genie

Send feedback

Last updated on **Apr 30, 2026**

Manage clean rooms

This page describes how to manage clean rooms, including how to:

- Update a clean room owner and comment
- Add, remove, and edit data assets and notebooks
- Grant access to a clean room
- Monitor clean room activity
- Delete a clean room

These tasks can be performed by all collaborators in a clean room.

Before you begin

Managing clean rooms requires the following permissions, depending on the task:

- To view a clean room in the clean room list or to view clean room details, you must be the owner of the clean room, a metastore admin, or have one of the following privileges on the clean room: `MANAGE`, `MODIFY CLEAN ROOM`, `EXECUTE CLEAN ROOM TASK`, or `BROWSE`.
- To update the owner of a clean room, you must be the owner of the clean room, a metastore admin, or have the `MANAGE` privilege on the clean room. The creator is automatically assigned as the owner of the clean room in their Databricks account. The collaborator organization's metastore admin is automatically assigned ownership of the clean room in their Databricks account. You can transfer ownership. See [Manage object ownership](#).
- To add and remove data assets and notebooks in a clean room and to update a comment, you must be the owner of the clean room or have the `MODIFY CLEAN ROOM` privilege on the clean room.
- To grant access to a clean room, you must be the owner, a metastore admin, or have the `MANAGE` privilege on the clean room.

- To delete a clean room, you must be the owner or have the **MANAGE** privilege on the clean room.

Update a clean room

1. In your Databricks workspace, click  **Catalog**.

2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. Select the clean room from the list.

4. Make any of the following updates:

- **Edit comment:** Click the  edit icon next to the comment. Comment changes apply only to the clean room securable in your Unity Catalog metastore. It does not get propagated to other collaborators.
- **Transfer owner:** Click the  edit icon next to the **Owner** name.
- **Add tables, volumes, and views:** See [Step 3. Add data assets and notebooks to the clean room.](#)
- **Remove tables, volumes, and views:** Click the  kebab menu (also known as the three-dot menu) on the asset row and select **Remove data asset**.
- **Add notebooks:** See [Step 3. Add data assets and notebooks to the clean room.](#)
- **Update notebooks:** Click the  kebab menu on the notebook row and select **Update notebook**. In the dialog, browse for and select the updated notebook.

Make any updates directly in the notebook before you select it in the Clean Rooms interface.

- **Remove notebooks:** Click the  kebab menu on the notebook row and select **Delete notebook**.

Grant access to a clean room

Here's a breakdown of clean room ownership and privilege assignment:

- The clean room owner has all privileges on the clean room.
- The clean room owner, metastore admin, and users with manage privilege can grant other principals access to the clean room.
- The creator is automatically assigned as the owner of the clean room in their Databricks account.
- The collaborator organization's metastore admin is automatically assigned ownership of the clean room in their Databricks account.
- You can transfer ownership. See [Manage object ownership](#).

The owner, metastore admin, and users with the `MANAGE` privilege can grant the following privileges on a clean room:

- `BROWSE`, which lets you list the clean room and view details like comments, owner, assets, and run history.
- `MODIFY CLEAN ROOM`, which lets you do everything that `BROWSE` grants, plus add and remove data assets, add and remove notebooks, and update comments.
- `EXECUTE CLEAN ROOM TASK`, which lets you do everything that `BROWSE` grants, plus run notebooks in a clean room.
- `MANAGE`, which lets you do everything that `BROWSE` grants, plus grant permissions to yourself or other users, and change the owner.

To grant a principal privileges on a clean room:

1. In your Databricks workspace, click  **Catalog**.
2. At the top of the **Catalog** pane, click the  gear icon and select **Clean Rooms**.

Alternatively, from the **Quick access** page, click the **Clean Rooms >** button.

3. Select the clean room from the list.
4. Go to the **Permissions** tab.

5. Select the user, group, or service principal, click **Grant**, and select the privileges you want to grant.

For more information about granting privileges, see [Manage privileges in Unity Catalog](#).

Monitor clean room notebook runs

All notebook runs from all collaborators are logged on the **Runs** tab in the Clean Rooms UI. You can filter runs by status and collaborator who ran the task.

1. In your Databricks workspace, click  **Catalog**.

2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. Select the clean room from the list.

4. Go to the **Runs** tab.

You can also view runs for a specific notebook by selecting the notebook in the clean room UI and going to the **Runs** tab on the notebook details page.

Monitor clean room logs

The `clean_room_events` system table logs all clean-room related events in your Databricks metastore and enables you to view all actions that you and other collaborators took on your clean rooms.

To learn how to use this system table to gain insight into clean room activities, see [Clean room events system table reference](#).

Delete a clean room

When you delete a clean room, collaborators are unable to use it, but it appears on their list of clean rooms until they delete the clean room object on their side. When

you delete a clean room, running tasks are canceled immediately.

1. In your Databricks workspace, click  **Catalog**.

2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. Select the clean room from the list.

4. Click the  kebab menu on the upper right corner of the page.

5. Select **Delete**.

6. Confirm that you want to delete the clean room, and click **Delete**.

Troubleshoot clean room issues

You may encounter the following error messages or issues when you work with clean rooms.

Other collaborators have left the clean room, so the clean room is no longer usable. Please delete this clean room securable.

In a no-trust clean room, any collaborator can delete the central clean room. This error message means that the other collaborator deleted the central clean room, and your reference to the central clean room is no longer valid. You can only delete the clean room object in your Databricks workspace. You cannot use it.

Collaborator X already has a clean room named Y

No collaborator can change the clean room name. This ensures that all collaborators can reference the clean room name without ambiguity. This error means that the other collaborator already has a clean room with the same name as the one you chose. Choose another clean room name.

Please accept the Serverless terms of service before using Clean Rooms

Serverless compute is required for central clean rooms. You do not need to enable serverless compute in your own workspace to use clean rooms. But an account admin needs to accept the serverless compute terms of service. If you need assistance, contact your Databricks representative.

Is this page

as this page helpful?

☐ Yes

☐ No

Last updated on **Dec 18, 2025**

Work with Databricks clean rooms as an invited collaborator

This page describes how to join and work in a Databricks clean room as an invited collaborator. A clean room can have ten total collaborators.

Databricks Clean Rooms provides a secure and privacy-protecting environment where multiple parties can work together on sensitive enterprise data without direct access to each other's data. For more information, see [What is Databricks Clean Rooms?](#).

When another collaborator creates a clean room so that you can work with them on a shared data project, there are two tasks that someone in your organization must perform:

1. Find and send the clean room creator your *Clean Room sharing identifier*.
2. When the clean room appears in your Unity Catalog metastore, grant users in your organization access to it.

This article describes those tasks, along with instructions for managing and working in a clean room.

Find your sharing identifier

When a collaborator creates a clean room that they want to share with you, they will request your Clean Room sharing identifier. The sharing identifier is a string that consists of your global metastore ID + your workspace ID + your username (email address).


Permissions required: To get the sharing identifier, you must have the `CREATE CLEAN ROOM` privilege or be a metastore admin.

To get the sharing identifier:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. On the **Clean Rooms** page, click your organization name in the upper right and select **Copy Clean Room sharing identifier** from the menu.



4. Send the sharing identifier to the clean room creator.

View, edit, grant access to, monitor, and delete a clean room

To learn how to view, edit, grant access to, monitor audit logs for, and delete a clean room, see [Manage clean rooms](#).

Add data assets and notebooks to a clean room that was shared with you

To learn how to add tables, volumes, views, and notebooks to a clean room, see [Step 3. Add data assets and notebooks to the clean room](#). The process is the same, whether you are the clean room creator or an invited collaborator.

Run notebooks in a clean room

To learn how to run notebooks in a clean room, see [Run notebooks in clean rooms](#).

Read output tables in a clean room

Output tables are temporary read-only tables generated by a notebook run and shared to the notebook runner's Unity Catalog metastore. If the notebook creates output tables, you can read those tables in your Databricks workspace, subject to the usual Unity Catalog privilege grants, and incorporate them into workflows. See [Create and work with output tables in Databricks Clean Rooms](#).

Is this page

as this page helpful?

☐ Yes

☐ No