

Last updated on **Apr 21, 2026**

What is Databricks Clean Rooms?

This page introduces Clean Rooms, a Databricks feature that uses Delta Sharing and serverless compute to provide a secure and privacy-protecting environment where multiple parties can work together on sensitive enterprise data without direct access to each other's data.

Requirements

To be eligible to use clean rooms, you must have:

- A workspace that is enabled for serverless compute. See [Serverless compute requirements](#).
- A workspace that is enabled for Unity Catalog. See [Enable a workspace for Unity Catalog](#).
- Delta Sharing enabled for your Unity Catalog metastore. See [Enable Delta Sharing on a metastore](#).

How does Clean Rooms work?


When you create a clean room, you create the following:

- A securable clean room object in your Unity Catalog metastore.
- The “central” clean room, which is an isolated ephemeral environment managed by Databricks.
- A securable clean room object in your collaborator's Unity Catalog metastore.

Tables, volumes (non-tabular data), views, and notebooks that either co-own or are shared in the clean room are shared with the central clean room only, using Delta Sharing.  Ask Genie

Sharing.

Collaborators cannot see the data in other collaborators' tables, views, or volumes, but they can see column names and column types, and they can run approved notebook code that operates over the data assets. The notebook code runs in the central clean room. Notebooks can also generate [output tables](#) that let your collaborator temporarily save read-only output to their Unity Catalog metastore so they can work with it in their workspaces.



How does Clean Rooms ensure a no-trust environment?

The Databricks Clean Rooms model is “no-trust.” All collaborators in a no-trust clean room have equal privileges, including the creator of the clean room. Clean Rooms is designed to prevent the running of unauthorized code and the unauthorized sharing of data. This trust is enforced explicitly by requiring all collaborators to approve a notebook before it can be run. While you can upload a notebook to yourself, it must be approved by the other collaborators before running. See [Approve a notebook in a clean room](#).

Additional safeguards or restrictions

The following safeguards are in place in addition to the explicit notebook approval process mentioned above:

- After a clean room is created, it is locked down to prevent new collaborators from joining the clean room.
- If any collaborator deletes the clean room, the central clean room is void and no clean room tasks can be run by any user.
- Each clean room is limited to ten collaborators.
- You cannot rename the clean room.

The clean room name must be unique in every collaborator's metastore, so that all collaborators can refer to the same clean room unambiguously.

- When a collaborator adds comments to a clean room securable in their workspace, these comments are not propagated to other collaborators.

What is shared with other collaborators?

- Clean room name.
- Cloud and region of the central clean room.
- Your organization name (which can be any name you choose).
- Your clean room sharing identifier (global metastore ID + workspace ID + user email address).
- Aliases of shared tables, views, or volumes.
- Column metadata (column name or alias and type).
- Notebooks (read-only).
- [Output tables](#) (read-only, temporary).
- Clean room events system table.
- Run history, including:
 - The name of the notebook that is being run
 - Collaborator that ran the notebook (not user).
 - The state of the notebook run.
 - The start time of the notebook run.

What is shared with the central clean room?

- Everything that is listed in the previous section.
- Read-only tables, volumes, views, and notebooks.

Tables, views, and volumes are registered in the central clean room's metastore with any supplied aliases. Data assets are shared throughout the lifecycle of the clean room.

Clean Rooms FAQ

The following are frequently asked questions about clean rooms.

How is my data managed in a clean room?

The central clean room is managed by Databricks. In the central clean room:

- Neither party has admin privileges.
- Only metadata is visible to all parties.
- Each party can add data to the central clean room.
- Clean rooms use Delta Sharing to share data securely to the clean room, but not between participants. See [What is Delta Sharing?](#).

How is my data kept private?

Central clean rooms run in an isolated, Databricks-managed serverless compute plane hosted in a cloud provider region that the clean room creator chooses.

Clean rooms provide:

- **Code approval:** The clean room creator and collaborators can share tables and volumes with the central clean room but can only run notebooks uploaded by the other party. You can review the code added by the other party before approving. If you run a notebook added by another party, you implicitly approve the code.
- **Version control:** Clean rooms notebooks have version control to ensure that all parties can only run fully-approved notebooks. Only the most recent version of a

notebook can be run. You can use the clean rooms system table to see which version of the notebook was run and monitor any changes made.

- **Restricted access:** When you create a clean room, you can use serverless egress control to manage outbound network connections. If you restrict access from your clean room, access to unauthorized storage is blocked. See [What is serverless egress control?](#)

To learn more about security and the serverless compute plane, see [Serverless compute plane networking](#).

How are actions recorded?

Clean room actions taken by you or your collaborators are recorded in the clean room events system table. These records include detailed metadata about the specific action taken. See [Clean room events system table reference](#).

Clean room actions are also recorded in your account's audit log under the service `clean-room`. See [Clean Rooms events](#).

When a collaborator modifies permissions on a clean room securable within their metastore, these changes are recorded in their audit logs under the `unityCatalog` service.

How does billing work?

To learn more about Databricks Clean Rooms pricing, see [link](#).

Limitations

The following limitations apply:

- No service credential Scala libraries included in the required Databricks Runtime version.
- If you use default storage tables in a clean room, Databricks recommends enabling the **Delta Sharing for Default Storage – Expanded Access** feature ([Beta](#)) in your account console.

Resource quotas

Databricks enforces resource quotas on all Clean Room securable objects. These quotas are listed in [Resource limits](#). If you expect to exceed these resource limits, contact your Databricks account team.

You can monitor your quota usage using the Unity Catalog resource quotas APIs. See [Monitor your usage of Unity Catalog resource quotas](#).

Get started

- [Create clean rooms](#)
- [Run notebooks in clean rooms](#)
- [Manage clean rooms](#)
- [Work with Databricks clean rooms as an invited collaborator](#)

Is this page

as this page helpful?

☐ Yes

☐ No

☐ Send feedback

Last updated on **Apr 21, 2026**

Create clean rooms

This page describes how to create a Databricks clean room using the UI. A clean room is a secure environment for collaborative data analysis.

Key features and limitations:

- **Secure collaboration:** Clean rooms enable multiple parties to work together on sensitive enterprise data without direct access to each other's raw data.
- **Collaborator capacity:** A clean room can have ten parties, including the creator and up to nine other collaborators.
- **Metastore limit:** Databricks enforces resource quotas on all clean room securable objects. See [Resource limits](#).

To use the REST API, see [Create a clean room](#).

Before you begin

The privileges needed to use clean rooms vary depending on the task:

Task	Required privileges
View a clean room	Must be the owner of the clean room, or have one of the following privileges on the clean room: <code>MANAGE</code> , <code>MODIFY CLEAN ROOM</code> , <code>EXECUTE CLEAN ROOM TASK</code> , or <code>BROWSE</code> .
Update the owner of a clean room	Must be the owner of the clean room, or have the <code>MANAGE</code> privilege on the clean room.
Add or remove data assets in	Must be the owner of the clean room or have the <code>MODIFY CLEAN ROOM</code> privilege on the clean room. If you are not the c

Task	Required privileges
a clean room	clean room, you and the clean room owner must have <code>SELECT</code> on any table or view and <code>READ VOLUME</code> on any volume that you add, along with <code>USE CATALOG</code> and <code>USE SCHEMA</code> on the parent catalog and schema.
Add or remove notebooks in a clean room	For the notebook's uploader: • If they are the designated runner of the notebook, they must have the <code>EXECUTE CLEAN ROOM TASK</code> or <code>MODIFY CLEAN ROOM</code> privileges. • If the notebook is runnable by a collaborator, the uploader must have the <code>MODIFY CLEAN ROOM</code> privilege.
Update a comment in a clean room	Must be the owner of the clean room or have the <code>MODIFY CLEAN ROOM</code> privilege on the clean room.
Grant access to a clean room	Must be the owner, or have the <code>MANAGE</code> privilege on the clean room.
Delete a clean room	Must be the owner or have the <code>MANAGE</code> privilege on the clean room.

Beyond the task-specific privileges, when a clean room is shared, the email address attached to the sharing identifier automatically becomes the owner in your collaborator's organization. See [Step 1. Request the collaborator's sharing identifier](#).

To learn about permission requirements for updating clean rooms and running tasks (notebooks) in clean rooms, see [Manage clean rooms](#) and [Run notebooks in clean rooms](#).

NOTE

The central clean room can have a maximum of two other non-AWS regions among its collaborators.

 Ask Genie

Step 1. Request the collaborator's sharing identifier

Before you can create a clean room, you must have the Clean Room sharing identifier of the organizations you will be collaborating with. The sharing identifier is a string that consists of the organization's global metastore ID + workspace ID + the contact's username (email address). Your collaborators can be in any cloud or region.

Reach out to your collaborators to request their sharing identifier. They can get the sharing identifier using the instructions in [Find your sharing identifier](#).

Step 2. Create a clean room

To create a clean room, you must use Catalog Explorer.

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.
3. Click **Create Clean Room**.
4. On the **Create Clean Room** page, enter a user-friendly name for the clean room.

The name cannot use spaces, periods, or forward slashes (/).

Once it's saved, the clean room name cannot be changed. Use a name that potential collaborators will find useful and descriptive.

5. Select the cloud provider and region where the central clean room will be created.

The cloud provider must match your current workspace, but the region can be different. Consider your organization's data residency or other policies when you make your selection.

6. Each clean room can have up to ten collaborators. Enter the **Clean Room sharing identifier** for each collaborator. See [Step 1. Request the collaborator's sharing identifier](#).

Create clean room

[Send feedback](#)

General

Set up the clean room

Clean Room name*

Region*

This will be the cloud and region for your secure central clean room.

us-west-2

Collaborator details

Collaborators cannot be modified after the clean room is created.

Clean Room sharing identifier* ⓘ

aws-eu-west-1:abcde

+ Add collaborator

You can test your clean room before full deployment by using either your sharing identifier or the identifier of another user in your current metastore. Doing so creates two clean rooms in your current metastore. For example, if you create a clean room titled `test_clean_room`, a second clean room named `test_clean_room_collaborator` also appears. Running notebooks with a collaborator in the same metastore functions the same as with an external collaborator. See [Run notebooks in clean rooms](#).

7. Make a note of the catalog names assigned to you and your collaborators.

All data assets added to the clean room will appear under that catalog in the central clean room, and can be referenced using that catalog in the Unity Catalog three-level namespace (`<catalog>.<schema>.<table-etc>`).

8. Select the network access policy type. This cannot be changed after the clean room is created.

- **Full Access:** Unrestricted outbound internet access.
- **Restricted Access:** This limits outbound access to internet destinations that you specify. See [Network policy overview](#) and [Manage network policies for serverless egress control](#).

Security

Security settings cannot be modified after the clean room is created.

Network access

Control which network destinations can be accessed from your clean room. Access to the cloud storage destinations backing your clean room assets will be automatically allowed.

☐ Full internet access
All internet destinations will be accessible from your clean room.

☒ Restricted internet access
Only internet destinations added to this policy will be accessible from your clean room.

+ Add destination

Restricted access can delay asset availability for up to ten minutes.

After you create the clean room, you can view the network access policy in the **Security** tab.

1. Click **Create Clean Room**.

If your current workspace is set to the HIPAA compliance security profile, then when you create a clean room, that setting is applied to the central clean room.

Collaborators must access the clean room from a workspace with the same security profile. See [Compliance security profile](#).

When creating a clean room, its security profile is inherited from your current workspace settings. This applies if your workspace is configured for HIPAA or FedRAMP Moderate compliance. FedRAMP Moderate is supported only when the clean room region is set to `us-west-2` or `us-east-1`. All collaborators must access the clean room from a workspace configured with the identical security profile to ensure consistent compliance.

Step 3. Add data assets and notebooks to the clean room

Both the creator and the collaborators can add tables, volumes, views, and notebooks to the clean room.

NOTE

The following instructions assume you are returning to an already-created clean room to add assets. If you just created a clean room for the first time, a wizard walks you through adding data assets and notebooks. The actual UI for adding these assets is the same, regardless of whether you are guided by the wizard or not.

To add notebooks:

1. Click the **+ Add notebooks** button and browse for the notebook you want to add.

 Ask Genie

2. Name the notebook.

3. Select which collaborator can run the notebook. Select **You** to run the notebook yourself.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Add notebook

Notebooks will be uploaded as static files to the clean room.

Notebook
/Users/test.user@collaborator1.com/Analysis Notebook

Notebook name
Analysis Notebook Preview

Notebook runner
Collaborator 1 (You) ▼

- ✓ Collaborator 1 (You)
- Collaborator 2
- Collaborator 3
- Collaborator 4

You can optionally give the notebook an alternative **Notebook name**.

Notebooks that you share in clean rooms query data and run data analysis workloads on the tables, views, and volumes that you and any other collaborators have added to the clean room.

If you share a notebook that includes results, those results are shared with your collaborators.

You can use a notebook to create *output tables* that are temporarily shared to your collaborator's metastore when they run the notebook. See [Create and work with output tables in Databricks Clean Rooms](#).

To use a test dataset, download our [sample notebook](#).

❗ IMPORTANT

Any notebook references to tables, views, or volumes that were added to the clean room must use the catalog name assigned when the clean room was created ("creator" for data assets added by the clean room creator, and "collaborator" for data assets added by the invited collaborator). For

example, a table added by the creator could be named `creator.sales.california`.

Likewise, verify that the notebook uses any aliases assigned to data assets that were added to the clean room.

To add assets:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms >** button.
3. Find and click the name of the clean room you want to update.
4. Click **+ Add data assets** to add tables, volumes, or views.
5. Select the data assets you want to share and click **Add data assets**.

When you share a table, volume, or view, you can optionally add an alias. The alias name is the only name visible in the clean room.

When you share a table, you can optionally add partition clauses that enable you to share only part of the table. For details about how to use partitions to limit what you share, see [Specify table partitions to share](#).

 NOTE

If you use default storage tables in a clean room, Databricks recommends enabling the **Delta Sharing for Default Storage – Expanded Access** feature ([Beta](#)) in your account console.

 NOTE

Federated table sharing is in [Beta](#). See [Add foreign schemas or tables to a share](#).

[Data sharing](#)[Clean Rooms](#)

Run notebooks in clean rooms

Last updated on **Apr 30, 2026**

Run notebooks in clean rooms

This page describes how to run notebooks in clean rooms. Notebooks are the interface that collaborators use to run data analysis in collaboration.

To learn how to add a notebook to a clean room, see [Create clean rooms](#).

Before you begin

To manage and run a notebook in a clean room, you must:

[Ask Genie](#)

Task	Requirement	Required privileges	Description
Managing a notebook	No unique requirements	For the notebook's uploader: - If they are the designated runner, they must have <code>EXECUTE CLEAN ROOM TASK</code> or <code>MODIFY CLEAN ROOM</code> privileges. - If a collaborator is the designated runner, the uploader must have the <code>MODIFY CLEAN ROOM</code> privilege.	General management tasks for a notebook in the clean room.
Running a notebook	Every collaborator, except the uploader, must approve the notebook. A designated runner runs the notebook.	<code>EXECUTE CLEAN ROOM TASK</code>	If the runner of the notebook did not upload it, they must approve the notebook before it can be run. This explicit approval can be automated through the default auto-approval rule. See Auto-approval rules .
Approving or rejecting a	None	<code>MODIFY CLEAN ROOM</code>	Allows you to  Ask Genie approve or reject

Task	Requirement	Required privileges	Description
notebook			notebooks before they can be run.
Managing auto-approval rules	Auto-approvals can only be applied to notebooks authored by collaborators other than the runner.	Clean room owner or <code>MANAGE</code> privilege.	Controls the automatic approval of notebooks.
Collaborator capacity	A clean room can include up to 10 collaborators.	None	This includes the creator and up to 9 other collaborators.

The creator is automatically assigned as the owner of the clean room in their Databricks account. The collaborator organization's metastore admin is automatically assigned ownership of the clean room in their Databricks account. You can transfer ownership. See [Manage object ownership](#).

Uploading a notebook and designating runners

When a collaborator adds a notebook to the clean room:

- They are considered the uploader of that notebook in the clean room.
- When adding a notebook, the uploader must designate which collaborator is the designated runner of the notebook. The designated runner is permitted to run the notebook.
- There can only be one designated runner.
- You can assign yourself as the designated runner of the notebook.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Add notebook

Notebooks will be uploaded as static files to the clean room.

Notebook

/Users/test.user@collaborator1.com/Analysis Notebook 

Notebook name

Analysis Notebook Preview

Notebook runner

Collaborator 1 (You) ▼

- ✓ Collaborator 1 (You)
- Collaborator 2
- Collaborator 3
- Collaborator 4



See [Step 3. Add data assets and notebooks to the clean room](#)

Approve a notebook in a clean room

Every notebook requires approval from all collaborators, except the uploader, before executing it.

To check the review status of a notebook:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms >** button.
3. Select the clean room from the list.
4. Select the notebook you want to run.
5. Click the **People** icon on the right to expand the collaborator section of the notebook details page.
6. The **Reviews** section provides a clear overview of which collaborators have approved, rejected, or have yet to review the notebook.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Analysis Notebook

Preview Runs

```
%sql
select * from collaborator_1.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_2.ads_demo.audience_data;
```

```
%sql
select * from collaborator_3.audience_overlap.advertiser_customers;
```

Collaborators New

Permissions

Owner	Collaborator 2
Runner	Collaborator 4

Reviews 2/4 approved

- ⌵ **Collaborator 1 (You)**
Pending review
- ✅ **Collaborator 2**
Added on May 12, 2025, 03:55 PM
System-generated
- ⌵ **Collaborator 3**
Pending review
- ✅ **Collaborator 4**
Auto-approved on May 12, 2025, 03:55 PM

A difference view is available when a notebook has been changed from a previous version.

When running a notebook version, the following is true:

- You can only run the latest version of a notebook.
- You can only approve or reject the latest version of a notebook.
- Modifying a notebook by adding new content or changing the runner designation creates a new version, which resets the review state for all collaborators.

To access the difference view for a notebook:

1. Click the **Clock icon** on the right to expand the version history of the notebook details page.

Catalog Explorer > Clean Rooms > multi-collaborator-clean-room >

Analysis notebook 3

Preview Runs

You are viewing the latest version from Apr 18, 2025, 10:29 AM. [Exit](#)

```
%sql
select * from collaborator_2.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_1.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_1.ads_demo.audience_data;
```

```
%sql
select * from collaborator_2.ads_demo.audience_data;
```

```
%sql
select * from collaborator_3.audience_overlap.advertiser_customers;
```

```
%sql
select * from collaborator_3.audience_overlap.advertiser_customers;
```


Version history

- Apr 18, 2025, 10:29 AM
✅ You approved
- Apr 18, 2025, 09:58 AM
✅ You approved
- Apr 18, 2025, 07:11 AM
✅ You approved
- Apr 16, 2025, 03:43 PM
✅ You approved

Ask Genie

To approve or reject a notebook:

1. Click the **Review** button at the upper right.
2. Choose to **Approve** or **Reject** the notebook.



The screenshot shows a modal dialog box titled "Update review for 'Analysis Notebook'" with a close button (X) in the top right corner. The dialog is divided into three sections:

- 1. Review notebook content**: A text area with the instruction "Confirm that the use of any private libraries and notebook parameters are correct."
- 2. Review notebook runner**: A dropdown menu showing "Collaborator 4" with a document icon.
- 3. Add a comment (optional)**: A large text input field.

At the bottom of the dialog are two buttons: "Reject" with a red 'X' icon and "Approve" with a green checkmark icon.

You can update your review up to nine times, for a maximum of ten reviews per collaborator on each notebook version. However, you can always reject a notebook, even if the review limit has been reached.

Auto-approval rules


Collaborators can set up auto-approval rules for their clean room. With auto-approvals, the following is true:

- You can create auto-approval rules only for notebooks uploaded by other users, not your own. Self-authored notebooks do not require your own approval if you are the designated runner.
- In 2-person clean rooms, you can auto-approve notebooks authored by the other collaborator.
- In clean rooms with more than two collaborators, you can auto-approve notebooks authored by **Anyone** or a specific collaborator.

- Each auto-approval rule designates a single runner for the approved notebook.
- Auto-approval is the default for notebooks uploaded by another collaborator when you are the designated runner.
- Auto-approvals are optional for clean rooms where the designated runner is another collaborator.
- You can add, update, or remove auto-approvals at any time.

To manage auto-approvals:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.
3. On the **Clean Rooms** page click the **Gear icon**
4. Select **Configure auto-approvals**.



Catalog Explorer > Clean Rooms >

 **multi-collaborator-clean-room**

Overview Runs Permissions Security


Filter assets Managed by: All collaborators

Notebooks (query data and run analysis)

Name	Status	Managed by	Added on	
 Analysis notebook 1	2/4 approved	e2-dogfood-clean-room-us-west-2	Apr 23, 2025, 10:20 AM	 Pending
 Analysis notebook 2	Rejected	e2-dogfood-clean-room-us-west-2	May 01, 2025, 11:07 AM	 Rejected

Configure auto-approvals

5. Configure automatic notebook approval based on matching permission rules.



Run a notebook in a clean room

To run a notebook, check all necessary approvals are in place and that you are the designated runner. Then:

1. Under **Notebooks**, click the notebook to open it in preview mode.
2. Click the **Run** button.
 - You can only run notebooks for which you are the designated runner and which have been approved.
 - It is not necessary to manually approve notebooks that you authored yourself to run them.
 - You can reject a notebook you authored. You must re-approve it before it can be run.
 - You can manage auto-approvals for each notebook. See [Auto-approval rules](#).
3. (Optional) On the **Run notebook with parameters** dialog, click **+ Add** to pass parameter values to the notebook job task.
4. Review the notebook.
5. Click **Run**.
6. Click **See details** to view the progress of the run. Alternatively, you can view run progress by going to **Runs** on this page or by clicking **Jobs & Pipelines** in the workspace sidebar and going to the **Job runs** tab.

7. View the results of the notebook run. The notebook results appear after the run is complete. To view past runs, go to **Runs** and click the link in the **Start time** column.

Even when all collaborators approve a notebook, only the collaborator designated by the uploader as the runner can run it.

Notebook parameters

The following parameters are automatically passed into the clean room notebook at runtime:

- `cr_central_id`: The central clean room ID.
- `cr_runner_global_metastore_id`: The global metastore ID of the designated runner.
- `cr_runner_alias`: The collaborator alias of the designated runner.
- `cr_<alias>_input_catalog`: The catalog that stores data shared by a specific collaborator, where `alias` represents their clean room alias. For two-party clean rooms created in the UI, this defaults to creator or collaborator, but it is customizable via the API. You can use `cr_<alias>_input_catalog` parameter for local tests with sample tables.
- `cr_output_catalog`: Defines the catalog where you create new output tables. Used in combination with `cr_output_schema`. See [Create an output table](#).
- `cr_output_schema`: Defines the schema where you create output tables. Used in combination with `cr_output_catalog`. See [Create an output table](#).

The above notebook parameters are automatically available as widget values during execution.

- **Python cells:** Reference them using `dbutils.widgets.get`. For example: `dbutils.widgets.get("cr_central_id")`.
- **SQL cells:** Access them with `select :name`. For example: `select :cr_central_id`.

Share notebook output using output tables

Output tables are temporary read-only tables generated by a notebook run and shared to the notebook runner's metastore. If the notebook creates an output table, the notebook runner can access it in an *output catalog* and share it with other users in their workspace. See [Create and work with output tables in Databricks Clean Rooms](#).

Use Lakeflow Jobs to run clean room notebooks

You can build complex, recurring workflows around your clean room assets using Databricks workflows. For example, you can create a recurring workflow that runs a clean room notebook task, then runs a task that immediately updates a report based on the clean room's output.

The following features facilitate complex clean room workflows:

- **Clean Room notebook task type:** Directly select and run a clean room notebook as a dedicated workflows task. See [Run notebooks in clean rooms](#).
- **Databricks Supplied output values:** All clean room notebooks tasks make the `{{tasks.<your_task_name>.output.catalog_name}}` and `{{tasks.<your_task_name>.output.schema_name}}` dynamic value references available to all downstream tasks. These values can be used for setting up workflows where a task is automatically passed the path to an upstream clean room notebook tasks's output schema. See [Supported value references](#)
 - `output.catalog_name` is automatically populated with the clean room's output catalog name.
 - `output.schema_name` is automatically populated with the clean room notebook task's dynamically-generated output schema name.
- **Lakeflow Jobs:** Use **Task values** that pass job parameter values to clean room notebooks or capture clean room notebook output and pass that output to other workflow tasks. See [Task values](#) to pass information between tasks.
- **Task values:** Like regular notebook tasks, clean room notebook tasks can set **Task values** that are passed to downstream workflow tasks. Use the syntax: `dbutils.jobs.taskValues.set(key="key", value="value")`. See more about [Task values](#).

Is this page

Was this page helpful?

☐ Yes	☐ No
Send feedback	

Last updated on **Dec 18, 2025**

Create and work with output tables in Databricks Clean Rooms

This page introduces output tables, which are temporary read-only tables generated by a notebook run and shared to the notebook runner's Unity Catalog metastore. This article describes how to use a notebook to create output tables and how runners can read these output tables in their Unity Catalog metastore.

Overview of output tables

Output tables let you temporarily save the output of notebooks that are run in a clean room to an output catalog in your Unity Catalog metastore, where you can make the data available to members of your team who don't have the ability to run the notebooks themselves. You can also use Lakeflow Jobs to run notebooks and perform tasks on output tables. Combined with the Clean Room notebook task type and support for task values, output tables let you create complex workflows that depend on Clean room notebooks.

Output tables are read-only.

Only the specific principal (user, group, or service principal) who runs the notebook has default read access to the output table. There is no write access. A metastore admin can grant read access to other principals in their Databricks account, using standard Unity Catalog privileges.

Output tables are stored for 30 days in the central clean room's default storage location and shared to the runner's metastore using Delta Sharing. If you want to keep an output table for more than 30 days, you must copy it to local storage.

Each notebook run creates a new schema in the output catalog. New runs cannot append an existing output table.

NOTE

Collaborators in Databricks on all three clouds—AWS, Azure, and Google Cloud—can share notebooks that create output tables and can read output tables that are generated when they run shared notebooks.

Create an output table

To create an output table, use the parameters `cr_output_catalog` and `cr_output_schema` in the three-part table namespace. Each run of the notebook produces a new schema.

In the following example, the notebook cell creates an output table called `overlapping_users` in the runner's output catalog that lists the users whose email address shows up in both the `collaborator.advertiser.profiles` and `creator.publisher.profiles` tables.

SQL

```
CREATE TABLE identifier(:cr_output_catalog || '.' || :cr_output_schema ||  
'overlapping_users') AS  
SELECT collab_profiles.*  
FROM collaborator.advertiser.profiles AS collab_profiles  
JOIN creator.publisher.profiles AS creator_profiles  
ON collab_profiles.email = creator_profiles.email
```

Read an output table

Output tables appear in a shared catalog in the notebook runner's metastore. In the Catalog Explorer **Catalog** pane, they appear in the **Shared** catalogs list.

Reading an output table is like reading any other table in Unity Catalog. You must have `SELECT` on the table, `USE CATALOG` on the shared output catalog, and `USE SCHEMA` on  Ask Genie

the automatically-generated schema. The user who ran the notebook that created the table has these permissions by default.

NOTE

Deleting a clean room removes all output tables and historical data from the output catalog.

Before you begin

This section describes cloud, configuration, and compute requirements for reading output tables.

Shared output catalog requirement

Before you can read output tables, a user must create the catalog that holds them. You only need to do this once per clean room. The owner of the clean room has permission to read and manage the output catalog by default.

Permissions required: EXECUTE CLEAN ROOM TASK

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.
3. Select the clean room from the list.
4. In the right pane, under **Output**, click **Create catalog**.
5. Enter an **Output catalog name** or accept the default, which is `<clean-room-name>_output`.

The output catalog appears in the list of **Shared** catalogs in your Catalog Explorer **Catalog** pane. Each clean room that you participate in can have one shared output catalog in your metastore.

Compute requirements

Queries on output tables require serverless compute. See [Connect to serverless compute](#).

Permissions required to read an output table

The user who ran the notebook that created the output table and the owner of the clean room have permission to read and manage the output table by default. All other users must have the following permissions granted to them:

- `SELECT` on the table
- `USE CATALOG` on the output catalog
- `USE SCHEMA` on the output schema

Run the notebook

To generate shared output tables in your output catalog, a user with access to the clean room must run the notebook. See [Run notebooks in clean rooms](#). Each notebook run creates a new output schema and table.

TIP

You can use Lakeflow Jobs to run notebooks and perform tasks on output tables, enabling complex workflows. See [Use Lakeflow Jobs to run clean room notebooks](#).

Find and view an output table

The user who runs the notebook that creates the output table can find a link to the output table on the notebook run history and run details pages in the **Clean Rooms** UI. In both cases, the link is in the **Output schema** field. See [Monitor clean room notebook runs](#).

Run history:

Catalog Explorer > Clean Rooms >						
 test-clean-room						
Overview Runs Permissions						
Notebook ▾ Run status ▾ Run by ▾ 1 job run						
Status	Start time	Run by	Notebook	Duration	Output schema	Expiration
	Sep 18, 2024, 03:08 PM	acme_co (You)	Output Tables - SQL	2m 38s	output_schema_57316933	Expires in 21 days

Run details:

Workflows > Runs > Run 251799029024345 >

CleanRoomsNotebookTask run Succeeded

Edit task Repair

Output

✓ 46.24 seconds 1: Create Output Table

```
%sql
CREATE TABLE identifier(
  :cr_output_catalog ||
  '.' ||
  :cr_output_schema ||
  '.overlapping_users'
) AS
SELECT collab_profiles.*
FROM collaborator.advertiser.profiles AS collab_profiles
JOIN creator.publisher.profiles AS creator_profiles
ON collab_profiles.email = creator_profiles.email;
```

_sqlidf: pyspark.sql.connect.dataframe.DataFrame = [num_affected_rows: long, num_inserted_rows: long]

Table

num_affected_rows

num_inserted_rows

No rows returned

0 rows | 46.24 seconds runtime

This result is stored as `_sqlidf` and can be used in other Python and SQL cells.

Task run

Details Metrics

Details

Job ID

672949637986455

Job run ID

349315244860569

Task run ID

573169339677688

Run as

Started

09/18/2024, 03:08:01 PM

Ended

09/18/2024, 03:10:44 PM

Duration

2m 43s

Queue duration

-

Status

Succeeded

Lineage

No lineage information for this job.
[Learn more](#)

View run events

View run libraries

Clean Room

Notebook

[test-clean-room/Output Tables - SQL](#)

Output schema

[test-clean-room_output/output_schema_...](#)

Expiration

Expires in 21 days

You can also find the output catalog in the list of **Shared** catalogs in your Catalog Explorer **Catalog** pane.

Limitations

In addition to the requirements listed in [Overview of output tables](#) and [Before you begin](#), output tables have the following limitations:

- Output tables are only supported when the clean room was created after the output table feature was released.
- Only tables are supported. Volumes and views, for example, are not.
- You are limited in the number of output tables each notebook can support. See [Resource limits](#).

Is this page
as this page helpful?

👍 Yes

👎 No

Ask Genie

Send feedback

Last updated on **Apr 30, 2026**

Manage clean rooms

This page describes how to manage clean rooms, including how to:

- Update a clean room owner and comment
- Add, remove, and edit data assets and notebooks
- Grant access to a clean room
- Monitor clean room activity
- Delete a clean room

These tasks can be performed by all collaborators in a clean room.

Before you begin

Managing clean rooms requires the following permissions, depending on the task:

- To view a clean room in the clean room list or to view clean room details, you must be the owner of the clean room, a metastore admin, or have one of the following privileges on the clean room: `MANAGE`, `MODIFY CLEAN ROOM`, `EXECUTE CLEAN ROOM TASK`, or `BROWSE`.
- To update the owner of a clean room, you must be the owner of the clean room, a metastore admin, or have the `MANAGE` privilege on the clean room. The creator is automatically assigned as the owner of the clean room in their Databricks account. The collaborator organization's metastore admin is automatically assigned ownership of the clean room in their Databricks account. You can transfer ownership. See [Manage object ownership](#).
- To add and remove data assets and notebooks in a clean room and to update a comment, you must be the owner of the clean room or have the `MODIFY CLEAN ROOM` privilege on the clean room.
- To grant access to a clean room, you must be the owner, a metastore admin, or have the `MANAGE` privilege on the clean room.

- To delete a clean room, you must be the owner or have the **MANAGE** privilege on the clean room.

Update a clean room

1. In your Databricks workspace, click  **Catalog**.

2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. Select the clean room from the list.

4. Make any of the following updates:

- **Edit comment:** Click the  edit icon next to the comment. Comment changes apply only to the clean room securable in your Unity Catalog metastore. It does not get propagated to other collaborators.
- **Transfer owner:** Click the  edit icon next to the **Owner** name.
- **Add tables, volumes, and views:** See [Step 3. Add data assets and notebooks to the clean room.](#)
- **Remove tables, volumes, and views:** Click the  kebab menu (also known as the three-dot menu) on the asset row and select **Remove data asset**.
- **Add notebooks:** See [Step 3. Add data assets and notebooks to the clean room.](#)
- **Update notebooks:** Click the  kebab menu on the notebook row and select **Update notebook**. In the dialog, browse for and select the updated notebook.

Make any updates directly in the notebook before you select it in the Clean Rooms interface.

- **Remove notebooks:** Click the  kebab menu on the notebook row and select **Delete notebook**.

Grant access to a clean room

Here's a breakdown of clean room ownership and privilege assignment:

- The clean room owner has all privileges on the clean room.
- The clean room owner, metastore admin, and users with manage privilege can grant other principals access to the clean room.
- The creator is automatically assigned as the owner of the clean room in their Databricks account.
- The collaborator organization's metastore admin is automatically assigned ownership of the clean room in their Databricks account.
- You can transfer ownership. See [Manage object ownership](#).

The owner, metastore admin, and users with the `MANAGE` privilege can grant the following privileges on a clean room:

- `BROWSE`, which lets you list the clean room and view details like comments, owner, assets, and run history.
- `MODIFY CLEAN ROOM`, which lets you do everything that `BROWSE` grants, plus add and remove data assets, add and remove notebooks, and update comments.
- `EXECUTE CLEAN ROOM TASK`, which lets you do everything that `BROWSE` grants, plus run notebooks in a clean room.
- `MANAGE`, which lets you do everything that `BROWSE` grants, plus grant permissions to yourself or other users, and change the owner.

To grant a principal privileges on a clean room:

1. In your Databricks workspace, click  **Catalog**.
2. At the top of the **Catalog** pane, click the  gear icon and select **Clean Rooms**.

Alternatively, from the **Quick access** page, click the **Clean Rooms >** button.

3. Select the clean room from the list.
4. Go to the **Permissions** tab.

5. Select the user, group, or service principal, click **Grant**, and select the privileges you want to grant.

For more information about granting privileges, see [Manage privileges in Unity Catalog](#).

Monitor clean room notebook runs

All notebook runs from all collaborators are logged on the **Runs** tab in the Clean Rooms UI. You can filter runs by status and collaborator who ran the task.

1. In your Databricks workspace, click  **Catalog**.

2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. Select the clean room from the list.

4. Go to the **Runs** tab.

You can also view runs for a specific notebook by selecting the notebook in the clean room UI and going to the **Runs** tab on the notebook details page.

Monitor clean room logs

The `clean_room_events` system table logs all clean-room related events in your Databricks metastore and enables you to view all actions that you and other collaborators took on your clean rooms.

To learn how to use this system table to gain insight into clean room activities, see [Clean room events system table reference](#).

Delete a clean room

When you delete a clean room, collaborators are unable to use it, but it appears on their list of clean rooms until they delete the clean room object on their side. When

you delete a clean room, running tasks are canceled immediately.

1. In your Databricks workspace, click  **Catalog**.

2. Click the **Clean Rooms >** button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. Select the clean room from the list.

4. Click the  kebab menu on the upper right corner of the page.

5. Select **Delete**.

6. Confirm that you want to delete the clean room, and click **Delete**.

Troubleshoot clean room issues

You may encounter the following error messages or issues when you work with clean rooms.

Other collaborators have left the clean room, so the clean room is no longer usable. Please delete this clean room securable.

In a no-trust clean room, any collaborator can delete the central clean room. This error message means that the other collaborator deleted the central clean room, and your reference to the central clean room is no longer valid. You can only delete the clean room object in your Databricks workspace. You cannot use it.

Collaborator X already has a clean room named Y

No collaborator can change the clean room name. This ensures that all collaborators can reference the clean room name without ambiguity. This error means that the other collaborator already has a clean room with the same name as the one you chose. Choose another clean room name.

Please accept the Serverless terms of service before using Clean Rooms

Serverless compute is required for central clean rooms. You do not need to enable serverless compute in your own workspace to use clean rooms. But an account admin needs to accept the serverless compute terms of service. If you need assistance, contact your Databricks representative.

Is this page

as this page helpful?

☐ Yes

☐ No

Last updated on **Dec 18, 2025**

Work with Databricks clean rooms as an invited collaborator

This page describes how to join and work in a Databricks clean room as an invited collaborator. A clean room can have ten total collaborators.

Databricks Clean Rooms provides a secure and privacy-protecting environment where multiple parties can work together on sensitive enterprise data without direct access to each other's data. For more information, see [What is Databricks Clean Rooms?](#).

When another collaborator creates a clean room so that you can work with them on a shared data project, there are two tasks that someone in your organization must perform:

1. Find and send the clean room creator your *Clean Room sharing identifier*.
2. When the clean room appears in your Unity Catalog metastore, grant users in your organization access to it.

This article describes those tasks, along with instructions for managing and working in a clean room.

Find your sharing identifier

When a collaborator creates a clean room that they want to share with you, they will request your Clean Room sharing identifier. The sharing identifier is a string that consists of your global metastore ID + your workspace ID + your username (email address).


Permissions required: To get the sharing identifier, you must have the `CREATE CLEAN ROOM` privilege or be a metastore admin.

To get the sharing identifier:

1. In your Databricks workspace, click  **Catalog**.
2. Click the **Clean Rooms** > button.

Alternatively, click the  gear icon at the top of the **Catalog** pane and select **Clean Rooms**.

3. On the **Clean Rooms** page, click your organization name in the upper right and select **Copy Clean Room sharing identifier** from the menu.



4. Send the sharing identifier to the clean room creator.

View, edit, grant access to, monitor, and delete a clean room

To learn how to view, edit, grant access to, monitor audit logs for, and delete a clean room, see [Manage clean rooms](#).

Add data assets and notebooks to a clean room that was shared with you

To learn how to add tables, volumes, views, and notebooks to a clean room, see [Step 3. Add data assets and notebooks to the clean room](#). The process is the same, whether you are the clean room creator or an invited collaborator.

Run notebooks in a clean room

To learn how to run notebooks in a clean room, see [Run notebooks in clean rooms](#).

Read output tables in a clean room

Output tables are temporary read-only tables generated by a notebook run and shared to the notebook runner's Unity Catalog metastore. If the notebook creates output tables, you can read those tables in your Databricks workspace, subject to the usual Unity Catalog privilege grants, and incorporate them into workflows. See [Create and work with output tables in Databricks Clean Rooms](#).

Is this page

as this page helpful?

☐ Yes

☐ No