

Last updated on **Mar 6, 2026**

Data engineering best practices

The following articles provide best practices for data engineering in Databricks.

- [Optimize join performance in Databricks](#)
- [Data modeling](#)
- [Configure RocksDB state store on Databricks](#)
- [Asynchronous state checkpointing for stateful queries](#)
- [Asynchronous progress tracking](#)
- [Production considerations for Structured Streaming](#)
- [Clean and validate data with batch or stream processing](#)
- [Observability in Databricks for jobs, Lakeflow Spark Declarative Pipelines, and Lakeflow Connect](#)
- [Fan-in and fan-out architecture in Lakeflow Spark Declarative Pipelines](#)
- [Best practices for Lakeflow Spark Declarative Pipelines](#)

For links to other best practices articles, including [CI/CD workflows best practices](#), see [Best practice](#).

Was this page helpful?

☐ Yes

☐ No

 Ask Genie